

Cadence: a patch architecture for continuing intelligence

Learning, memory and action through equilibrium repair

Bernhard Mueller

Draft of September 23, 2026

Abstract

Cadence investigates a general learning architecture built from bounded observer-like patches that retain state, exchange readings and repair inconsistent predictions. Overlap repair, equilibrium detuning and consequence readback organize a continuing loop: interpret experience, imagine actions, act, and revise learned relationships from actual outcomes. The proposed equilibrium represents the actor's current belief about its world, body and possible actions. Familiar behavior should follow a changing sequence of consistent states with inexpensive repair; new evidence updates that belief, while learning the underlying dynamics is a distinct operation. We hypothesize that recursive nonlinear composition and prediction grounded in demonstrations and interventions can support this organization. A second hypothesis is that metacognition is recursive self-observation: a bounded patch that reads and steers the rest of the brain must select what matters, and that selection is attention; a theorem shows that a tower of linear readbacks has depth zero, so a steering level must act on the equations below it. A ladder of such levels, from reflex to a robot that acts and speaks among people, orders the program's tasks. Every step is accepted by a demo: a page with the whole brain in the viewer, the world, an instrument strip, a three-way switch (the step on; off, the brain below it; the hand-designed control) and the person in the world, accepted when it works with the switch on and fails with it off. The first two rungs are earned under that rule. A governor patch that reads the brain's own surprise and returns its mode catches 0.961 of the laser dots on a sill awake 0.238 of the time at 15.3 patch moments per decision, where a brain always awake catches 0.962 at 71.3 and a brain that never wakes catches 0.249. A steering patch that reads a cortex's readback and returns the gain of each sense inside its repair predicts a dummy's next place with error 0.041 across three hidden phases, where fixed gains give 0.081 and the designed inverse-variance rule 0.062, weighing the eye at 1.8 when the voice is displaced and the ear at 1.8 in the dark. In every acceptance demo, selection over a genome beat the hand-set genome and a random search and pruned what did not pay. Controlled examples establish narrower capabilities: temporal learning, conditional response protection, action through acquired body models, exact past-boundary compression under a fixed Gaussian model, and acquisition in two phases on a finite sentence grammar, where records retain a corpus by day and slow weights learn their completions by night. Mathematical statements and independently checked experiments identify the assumptions and preserved quantities. Biological mechanisms motivate simplified abstractions; we do not claim brain equivalence or consciousness. Scalable world models, useful recursive self-reflection and reliable long-form creation are unestablished. Measured observations are kept separate from private imagined futures.

Contents

1	A continuing actor built from local repair	3
2	Shared patch operations and persistent state	4
3	Architecture and controlled evidence	12
4	Language against a next-word learner	18
5	The metacognition ladder	20
6	Biological abstractions and functional self-readback	27
7	Limits and the generalization test	28
8	Conclusion	28
A	Certified recurrence and local credit	31
B	Optional record algebra and capacity	34
C	Wiring and a declared cost model	36

D	Theorem map and the Lean library	36
E	Evidence map	38
F	Notation	41

1 A continuing actor built from local repair

The design question is the smallest persistent state and update rule that preserve useful experience. A capable continuing learner must acquire a body model, retain context and important relationships, imagine consequences, act, and correct itself from what actually happens. Cadence investigates whether these functions can share bounded patch operations instead of requiring a separate mechanism for each task. The individual is the continuing organization of activity, learned relationships, interfaces and body coupling; one scalar energy or a stable output is not a complete description of it.

Observer-Patch Holography (OPH) supplies the organizing constraint: bounded, self-reading patches exchange declared public readings and repair disagreement while protecting observations [42, 43]. Overlap consensus concerns compatible interface values, not identical internal states. Conditional normal-form results distinguish a consistent answer from an answer independent of the initial interior. Neither property establishes that the input evidence is true or that the learned model is useful. Section A.2 gives the exact neural specialization.

Three operations structure the proposal. *Repair* settles local state under held boundaries. *Detuning* changes relationships using local contrasts between equilibria perturbed by real teaching evidence. *Self-readback* makes proposed actions and predicted consequences available to the same continuing actor, whose next observation tests those predictions. Physical readback supplies new evidence; an imagined outcome does not. A preference can guide future action without rewriting the observed past.

The end of the program is a brain in a body that functions among people: it looks, acts and speaks, holds a plan, asks when unsure and says what it is doing; past that, brains that read their own development and one another. Section 2.7 states the hypothesis that carries the route, metacognition as recursive self-observation, with a theorem on what a self-reading level must do to be more than a reparameterization, and Section 5 orders the tasks between here and that end on a ladder of such levels. That section also states the rule by which a step is accepted, a demo that works with the step on and fails with it off against a hand-designed control, and reports the first two rungs as earned under it: a governor patch that reads the brain’s own surprise and returns its mode, and a steering patch that reads a cortex’s readback and returns the gain of each sense inside its repair.

These operations have established antecedents. Hopfield networks compute by collective relaxation [24]; equilibrium propagation derives local credit under symmetry and stable-branch assumptions [31, 54]; temporal predictive coding uses transition and readout residuals [35]. Integrated neural architectures and learned body models motivate the whole-actor goal [8, 12]. Cadence’s contribution is an explicit patch/evidence contract, mathematical boundaries on retained state, and controlled tests of which combinations actually support useful action. No superiority over conventional filtering, control or regression follows from sharing this description.

1.1 Scope of the evidence

The base library exposes a nonlinear temporal learner, explicit response constraints, private continuous-input planning, a fixed linear-body actor, and a record patch whose store takes an observation in one write by day and whose slow weights learn by night from the store’s own completions with the world closed (Section 4.1). These share residual-repair ideas but retain distinct state spaces and hypotheses. The temporal learner accepts time-varying observations, carries only target-free activity after detuning, and can repair action ports under its own acquired dynamics. Both control interfaces hold their body model fixed while separating actual records from hypothetical plans. An audio application must learn or supply its instrument relation. The belief patch carries a belief across moments by a learned transition, repairs it from evidence through a fixed number of iterations with the record store read inside the repair, and continues it in private imagination when evidence is withheld; record patches join through declared ports into one joint equilibrium (`cadence.record_ports`); a habit is a settling patch whose answer is its motor state; a steering patch is a patch of the same rule whose evidence is the readback of other patches. Sources, frozen protocols, controls and receipts identify which implementation supports each result.

The main evidence is presented in Sections 3, 4 and 5. Appendix A preserves the unique recurrence, record-interference and wiring theorems with their hypotheses. Appendix E retains the source-bound receipt map, including the application evidence that motivated these tests. The public proof package contains 169 audited statements, including the causal-path, quadratic-boundary and protected-path results. Numerical verification is separate from a formal refinement proof of the implementation.

The general behavioral criterion is a learned organization that changes its actual continuation correctly when memory, intention or readback is intervened upon, while preserving unrelated useful experience. In a creative task, it must also construct and evaluate alternatives beyond recalled examples. Musical composition is one demanding instance: heard context, retained intention, imagined alternatives and instrument consequences must influence the same continuing actor.

2 Shared patch operations and persistent state

2.1 Local state, learned relationships and ports

The primary library interface is **TemporalPatchNet**: learned transition/readout maps, carried free activity and explicit temporal boundaries support observed learning and private future repair [37]. **TemporalMemory** adds model-bound response constraints; **plan** changes declared control inputs on a private branch. The temporal model retains A, B, C , free activity, output precision and solver/revision metadata; protection additionally retains its bases, tolerance and parameter binding. Section 3.1 gives its residual energy. The optimizer and delivery-identifier state below belong to the supported graph specialization.

That graph specialization, **PatchNet**, wraps the nonlinear **Brain** and reciprocal **Learner**. It owns a continuing free neural state, learned synaptic efficacies and biases, and the learner’s optimizer state. Input and output ports name the coordinates through which observations and predictions cross its boundary. A caller supplies their meaning, the initial wiring and the teaching objective. Internal functional specialization is a learning question; declaring a region to be a workspace does not establish that it performs a useful workspace function.

For the graph specialization’s fast state, the local constraint is

$$v_i = \sum_j W_{ij} \text{act}(v_j) + c_i,$$

with the nonlinear activation and repair step specified in Section A. A patch reads its neighbors’ published values and its own boundary drive, then updates its discrepancy. Effective weights, rather than neuron labels or equal neighboring values, define the relationship being repaired. The software schedules these local dependencies centrally. An energy interpretation requires the effective reciprocal interactions and the activation assumptions of the learning derivation; arbitrary directed connectomes do not acquire that interpretation by sharing the interface.

Its retained computational state has the form

$$S_t = (v_t, \theta_t, o_t, P_t),$$

where v_t is free activity, θ_t the learned relationships, o_t the optimizer state and P_t a bounded optional record of delivered source identifiers. Fixed interface and solver configuration accompany this state. A complete reproducible system also accounts for any external random generator, body state and scheduler that affect subsequent observations. The optional identifier record provides duplicate-delivery bookkeeping; it is not an episodic content store.

Continuous learned weights express graded relationships. The minimal neural core retains activity and parameters without an attached content store or importance classifier. Its empirical obligation is to show which useful memories these states can acquire, preserve and revise; the absence of a separate store does not itself satisfy that obligation.

Ports between record patches form one joint equilibrium. A port carries a declared band of one record patch’s context, in the unit its records read it in, into another patch’s inputs in the same moment; the seam is a term of the joint energy and the adjoint crosses it. On two coupled symbol streams with one hidden cause, at 2,588 to 2,604 parameters for every arm, the joint port lifts the dependent stream’s held-out accuracy from 0.393 with no port to 0.544, level on the fill-in stretches with 0.598 for one wide patch over both streams, whose records key on both symbols at once; the same band read from the previous moment gives 0.463 and 0.468 with and without the gradient across the seam, so the same-moment settling round is what carries the current symbol; cutting the port from the trained brain drops the stream to 0.268, below the independent

Operation	State and evidence contract
create	Supplies a reciprocal initial graph and named input/output ports; these are inherited structure, not learned roles.
settle, read	Advances and reads target-free activity without changing learned parameters. Inputs are soft drives and receive recurrent feedback.
observe	Uses externally observed target coordinates and supplied gains; updates parameters only after all required phases converge. A rejected attempt can advance free activity.
imagine	Evolves a sequential private branch; neither live continuation state nor witnessed evidence is changed.
reset	Clears fast activity while retaining parameters, optimizer state and delivery identifiers.
save, load	Preserves continuation state; exact numerical continuation requires the same backend and precision.

Table 1: The supported graph specialization’s public interaction contract. These operations expose computation and evidence boundaries; they do not guarantee learned competence.

patch, since the patch came to rely on what crosses [R115]. On a ball that an eye and an ear watch, the port from the ear keeps the eye’s belief alive in the dark: an error of 0.230 with hand-designed ports against 0.280 for an eye alone and 0.329 with the ports cut at test. Selection over the topology as a genome, each port’s existence, band and width and the settling rounds, kept one port from the ear into the eye over a middle band of the ear’s timescales, at 0.169, and dropped the port into the ear, which cost channels and bought nothing the fitness saw. The ported ear names 0.03 of the coming bounces against 0.30 for the wide patch, a negative kept with its mechanism: a linear readout over a leaky context of a bump code decodes position with a lag, and the record patch extracts little velocity from the retina; the belief patch below is the building block that task asks for [R115].

2.2 An explicit temporal overlap

Carrying an initial state cannot preserve its information when complete settlement has one solution independent of that state. The graph specialization’s optional temporal boundary makes the previous free activity $\bar{s}$ an argument of the constraints themselves. With a supplied mask $a_i \in \{0, 1\}$ and strength $\kappa \geq 0$, its energy term is

$$E_{\bar{\theta}}^{\bar{s}}(s) = E_{\theta}(s) + \frac{\kappa}{2} \sum_i a_i (s_i - \bar{s}_i)^2.$$

The corresponding local recurrent equation adds $\kappa a_i (\bar{s}_i - s_i)$ to the drive. The same $\bar{s}$ is held throughout the free, positive and negative phases of one observed event. It is not advanced at each numerical iteration and is independent of the target’s perturbation strength and evidence gain. The next event uses the preceding free activity; a fresh or reset stream uses a zero boundary. An imagined branch advances only its private previous state.

The default strength is zero, and the default eligible coordinates are the inherited hidden ports. Positive strength is therefore a supplied persistence assumption. It introduces no additional fact store and reuses existing activity, but it changes the model’s equations and must be counted as a structural choice. The unmodified recurrence’s contraction certificate does not automatically certify this changed solver. In a single-minimum landscape, repeated temporal overlap can merely slow forgetting. It does not discover a protected subspace, choose important events or guarantee indefinite retention. The gradient estimate below is conditional on the held boundary; it does not differentiate through preceding events or solve long-delay temporal credit.

An exact scalar example distinguishes persistence from protection. For $a > 0$ and $\kappa \geq 0$, the energy $az^2/2 + \kappa(z - c)^2/2$ has unique minimizer qc , where $q = \kappa/(a + \kappa) \in [0, 1)$. Reusing that minimizer as the next boundary gives $q^n c$, tending to zero. These statements are machine-checked ¹. They describe a convex scalar system; they do not prove that every nonlinear neural basin washes out.

¹Lean: `repair_unique`, `iterate_closed_form`, `iterate_tends_to_zero`.

A readable trace is a different boundary. The library also provides **Trace**, which a caller connects explicitly to context ports. With movement weighting disabled, it records selected free activity h_t as

$$z_{t+1} = \rho z_t + (1 - \rho)h_t, \quad 0 \leq \rho < 1,$$

and makes z_t readable at declared context ports during the next event. The trace is held fixed across all numerical iterations and nudged phases, then advanced once from the free state. This supplies an explicit fading observation history; carrying a numerical initial state or tethering the immediately preceding activity implements a different operation.

For fixed z_t, h_t , the displayed update is the unique minimizer in u of

$$\frac{\rho}{2}(u - z_t)^2 + \frac{1 - \rho}{2}(u - h_t)^2.$$

Six machine-checked lemmas establish this scalar repair, interval preservation, decay of an initial difference as ρ^n under the same source sequence, and decay to zero under zero source ². Here the timescale and source are supplied. If the recurrent source itself depends on the trace, the same-source premise need not hold. These lemmas certify neither a joint energy for the full recurrent system nor its stability, learned importance or permanent memory [R80].

2.3 Readback and detuning as the learning operation

A witnessed target y and an observed-coordinate mask m define a port loss, for example

$$L(s; y, m) = \frac{1}{2} \sum_k m_k (s_k - y_k)^2.$$

Unobserved coordinates do not supply targets. At fixed parameters and fixed temporal boundary, the default centered learner obtains a free state and states under positive and negative perturbations $E_\theta^\pm \pm \beta L$. It then estimates a parameter derivative from the difference of the corresponding local energy derivatives:

$$g_e(\beta) = \frac{\partial_{\theta_e} E_\theta(s^+) - \partial_{\theta_e} E_\theta(s^-)}{2\beta}.$$

For the implemented pair interaction, this is an endpoint activity contrast. The current optimizer and parameter constraints determine the ensuing update. The gradient interpretation is imported from equilibrium propagation under its smooth stable-branch, symmetry and convergence assumptions [31, 54]; endpoint locality alone proves none of those assumptions. Finite nudges are biased, branches can differ in a multibasin system, and free-state convergence does not imply convergence of both nudged phases.

The revised interface checks the equation residual of every required phase and refuses a parameter update when a phase reaches its work cap without meeting tolerance. It reports convergence and work rather than treating small successive movement as sufficient evidence of settlement. This is a numerical acceptance contract. Outside the contractive subclass it is neither a bound on distance to a unique equilibrium nor a proof of local stability or correct task predictions.

Phase convergence also does not certify a useful parameter step. The temporal learner’s optional **backtrack** admission keeps the same centered contrast g_β and tests $\theta' = \theta - \eta g_\beta$ by an ordinary free replay from the original observed path boundary. It accepts only a finite, converged prediction satisfying

$$L(\theta') < L(\theta) - \epsilon_{\text{num}}, \quad L(\theta') \leq L(\theta) - 10^{-4} \eta \|g_\beta\|^2,$$

where L is the observed precision-weighted half mean squared error and ϵ_{num} is a declared floating-point floor. At most sixteen successively halved rates are tried. Rejection preserves parameters and their revision count; valid pre-update free activity can advance. Targets score the replay without driving its recurrence, and trial activity is private. This guards improvement on the current observation, not old skills, unseen data or the direction’s exact-gradient interpretation. The fixed-rate default is unchanged. Independent replay of the archived experimental guard checks 961 observed updates, including 576 initial proposals that require a

²Lean: `unique_minimizer`, `repair_mem_interval`, `same_source_history_difference`, `zero_source_tends_to_zero`.

smaller rate. A separate comparison verifies public-interface agreement on seven saved updates; it is not a full training replay through the public implementation [R95].

The distinction between activity repair and learning is essential. A free equilibrium answers the current constraints. A witnessed teaching perturbation changes persistent relationships only through the learning operation. An unobserved read does not supply a fabricated target. The caller’s loss, target and evidence selection therefore determine what counts as a useful correction; a harmful-outcome nudge is supplied valence, not autonomous discovery of importance.

2.4 Memory, imagination and evidence identity

Free activity carries forward to the next decision. Resetting fast activity leaves learned relationships intact, allowing a cold query to test parameter memory separately from a lingering response. Checkpoints preserve parameters, optimizer state, activity and bounded delivery bookkeeping so that restart is not confused with relearning. These contracts make retained state inspectable; task-specific retention must be measured.

Imagined alternatives settle from isolated copies of the free state. They cannot update the live state, learned parameters or witnessed-event bookkeeping. Their predictions can guide an external action, but agreement among simulated branches is not fresh evidence about the world. A branch that predicts a shock and an instrument that reads back a shock have different evidential roles. Learning from a generated target would require an explicitly different, disclosed training protocol.

This boundary is enforced by the branch operation’s lack of writes. The caller supplies observations to the learning interface; the library does not authenticate that a caller-provided target came from a physical measurement. Source identifiers authenticate neither that provenance nor the target’s truth.

An optional source identifier suppresses repeat delivery of the same observed payload while that identifier is retained, and rejects conflicting reuse of that identifier. Finite storage limits this guarantee. Distinct identifiers do not establish independent evidence; repeated reports from one sensor can share a cause. The core does not construct an authenticated causal graph, infer sensor reliability or decide whether a contradiction is noise or a changed world.

2.5 What an equilibrium can remember

Fast memory and parameter memory differ. At fixed drive and parameters, global contraction implies one equilibrium. Consequently complete settlement discards differences carried only by the initial activity. Warm starts save work but do not evade this implication. History can instead affect a persistent parameter, an explicit boundary value, or a basin choice outside global contraction. A multiple-basin system can preserve a cue under unchanged observations; learning changes its landscape and can destroy that cue while preserving the old cold input-response mapping.

Sufficiency concerns future behavior. Let h be a history, T a declared deterministic update, and r its required readout. The full future signature is

$$F(h)(u) = r(T^*(h, u)), \quad U(F, i)(u) = F(i :: u), \quad r_F(F) = F(\square),$$

where u ranges over finite input continuations. Histories with identical signatures admit the same shift update and present readout. Every exact memory with state-only updates and exact readout must distinguish histories with different signatures. These statements are machine-checked ³. This is a behavioral quotient, not a finite learned compression algorithm: the signature can be infinite and a callable representation may retain the whole history. A family of N behaviorally distinguishable histories requires at least N distinguishable memory states ⁴. Finite-precision parameters therefore have a capacity limit; counting real coordinates without precision does not remove it. The required task and continuation family determine which distinctions matter. The construction does not reveal arbitrary future importance to a learner [38].

³Lean: `future_signature_one_step, every_exact_dynamic_encoding_refines_future_signature`.

⁴Lean: `distinguishable_family_capacity_bound`.

A causal poset supplies dependence, not missing content. Two records with different values can have exactly the same causal order and maximal node identifiers. The corresponding topology-only encoding fails any query that distinguishes those values ⁵. A sufficient frontier must retain the required payloads or boundary summaries as well as causal context. Disjoint replacement writes commute under their update law; conflicting writes on one key do not ⁶. Provenance can therefore support safe delivery and revision, but a DAG alone provides neither storage semantics, learned importance nor confluence.

Normal form is not objective truth. A terminating, confluent repair relation gives a unique normal form from a starting state under its hypotheses. A boundary-determined answer additionally requires all relevant history to be represented in that boundary. Neither condition implies that noisy observations were interpreted correctly. In a two-world construction, the same available reports require latent answers zero and one. Any single estimate z then satisfies

$$z^2 + (z - 1)^2 \geq \frac{1}{2},$$

so at least one squared error is at least $1/4$ ⁷. This is an information bound, not a stochastic noise model. Immediate perfect revision and perfect immunity to an indistinguishable first outlier cannot both be required. Acquiring another informative observation is one possible response to that uncertainty.

Conditional basin retention. Let B be a supplied basin, let an allowed exit from B raise a reference energy by at least $m > 0$, and suppose each time-varying energy E_t differs from that same reference by at most ϵ at every state under consideration. If $2\epsilon < m$, an allowed step that does not increase E_t cannot leave B . Every finite accepted trace therefore stays in B , preserving any readout constant there ⁸. A per-step bound relative only to the preceding energy is insufficient for this conclusion. The theorem assumes the exit margin and allowed moves; it does not certify the numerical neural implementation’s basins, find their margins, or prove retention under arbitrary future learning.

Specialization needs asymmetry and useful consequences. A deterministic permutation-equivariant update preserves an invariant initial state under invariant inputs. Exchangeable owners with identical complete states and exposures cannot differentiate by that update alone ⁹. Random initial conditions, unequal experience or inherited asymmetry can remove this obstruction. They do not themselves establish useful division of labor. A symmetric task can have asymmetric optima. For a supplied mismatch r and connection price $c > 0$, minimizing $w(r^2 + c)$ over a freely adjustable weight $w \geq 0$ selects $w = 0$ ¹⁰. Retaining useful communication therefore needs a task constraint or consequence beyond cheaper agreement. The port search above is an instance: the port into the ear had a price in channels and no measured consequence, and selection dropped it [R115].

2.6 The equilibrium world model hypothesis

We call the intended organization a *Cadence equilibrium world model*; a patch net denotes the computational mechanism rather than an established level of competence.

The architectural hypothesis is that a repaired state can represent the actor’s current situation: persistent surroundings, its controlled body, available actions and resources, and uncertainty about their consequences. This is a distributed representation used by perception and action; it need not be expressed in words. Its adequacy is behavioral: histories that require different predictions or actions must retain distinguishable representations. A converged state can be a belief about an uncertain world without selecting one hypothetical outcome as fact. Numerical consistency alone establishes none of this semantic adequacy.

⁵Lean: `topology_only_is_insufficient, maximal_ids_only_are_insufficient.`

⁶Lean: `distinct_key_updates_commute, same_key_updates_do_not_commute.`

⁷Lean: `two_target_squared_errors, indistinguishable_noise_or_change_forces_error.`

⁸Lean: `finite_trace_preserves_basin, finite_trace_preserves_readout.`

⁹Lean: `equivariant_rollout_stays_fixed, homogeneous_owner_rollout.`

¹⁰Lean: `positive_cost_selects_zero.`

A changing trajectory can be consistent. Let z_t be the carried interpretation, a_t the executed action, o_t the observation and M_t retained evidence. A candidate realization separates prediction in physical time from repair within one decision:

$$\begin{aligned}\bar{z}_t &= T_\theta(z_{t-1}, a_{t-1}), & z_t^{(0)} &= \bar{z}_t, \\ z_t^{(k+1)} &= (1 - \alpha)z_t^{(k)} + \alpha F_\theta(z_t^{(k)}, o_t, \bar{z}_t, M_t), & 0 < \alpha &\leq 1.\end{aligned}\tag{1}$$

Here k counts inference steps while t counts experienced moments; T_θ predicts a continuation and F_θ assimilates evidence. These are proposed functional roles, not a claim that every implemented patch has this form. Reapplying F_θ at a fixed point does not advance physical time. A familiar game can instead trace a changing family of consistent states: the predicted continuation is close to the next adequate interpretation and a learned policy supplies appropriate action with little additional deliberation. Even predictable danger can require an immediate practiced response. An observer predicting a freely moving rock through $\hat{x}_{t+\Delta t} = x_t + v_t \Delta t$ illustrates continuation under a retained regularity: agreement permits inexpensive tracking, while a discrepancy requires interpretation. This local inertial analogy does not make physical momentum depend on the observer’s surprise, derive neural dynamics, or establish subjective experience.

Randomness revises belief without necessarily revising a law. A calibrated model of a random event predicts a distribution, so the transition prior in Equation (1) must represent relevant uncertainty. Observing one realization can substantially change the current interpretation and appropriate action even when that distribution was correct. Such an innovation does not by itself justify replacing the learned transition law. Individual observations can supply nonzero learning updates; calibration removes systematic prediction-error bias in expectation, rather than forbidding statistical learning. Repeated errors, assessed relative to predictive uncertainty, can instead support parameter revision. New sensory evidence, repair of a proposed action and the teaching perturbations $\pm \beta L$ therefore have different roles; an unexpected observation is not the numerical detuning parameter. The contraction results bound repair under their stated fixed-model assumptions. They do not imply that stochastic observations cease to arrive or that familiar behavior costs no computation.

Composition and grounded prediction. Two ingredients motivate Equation (1). First, scene, body, proposed action and retrieved context must interact nonlinearly, with one interpretation available to another inference step. For example, an additive predictor $\hat{y}(v, a) = f(v) + g(a)$ has an action difference $g(a_1) - g(a_2)$ independent of the scene v ; it cannot express a reversal of that difference between two scenes. Width or separate feature depth preserves this obstruction if the final combination stays additive. Second, observed consequences must train the representations, transition and retrieval that produced their predictions, including continuation from predicted states. These ingredients are hypotheses about useful representation learning, not sufficient conditions for intelligence. Nonlinear memory feedback requires its own convergence and credit analysis; the quadratic and contractive results do not transfer automatically.

Imitation supplies a starting point for understanding. Skilled trajectories can jointly teach a behavioral prior and the consequences of executed actions. Observation without action labels needs an additional identification mechanism: Video PreTraining uses an inverse-dynamics model trained on labeled data to infer actions in unlabeled video [5]. Copying expert actions alone does not establish the effects of alternatives, and a learner’s errors can take it outside the demonstrated distribution; DAgger addresses this distribution shift through expert supervision on learner-induced states [50]. A continuing actor should combine demonstrations with witnessed interventions and recovery experience, keeping demonstrator actions distinct from its own. Improvement beyond the teacher requires a better supported predictive or decision signal, not repetition alone.

Internal imagination and cheap familiar action. A private branch may advance T_θ under proposed actions and evaluate its predicted consequences without rendering a full sensory image at each step. Dreamer and MuZero provide precedents for learning useful behavior through learned-model imagination and planning [22, 56]; they do not establish the proposed equilibrium mechanism. A fully internal branch receives no future body observations, rewards or terminal flags from the real trajectory. Its benefit requires action comparisons against subsequent evidence and a declared computation budget. Grounded rehearsal can then train a faster

initialization or habitual policy from effective deliberation. The test is equal or better behavior with less additional repair or planning on unseen familiar situations, while retaining recovery from perturbations. Low energy in a dream is insufficient: an incorrect model can be internally consistent. Useful uncertainty and long-horizon imagination are unestablished general capabilities of Cadence; the private continuation and the reduction in decision cost are measured below on small worlds.

The belief patch carries a belief under a screen. The tested form of Equation (1) is the belief patch: a learned transition carries the belief to the next moment, evidence repairs it through two iterations of one map with the record store read inside the repair, and when evidence is withheld the transition alone continues it. The shell game measures the continuation: three identical cups, a ball under one, swaps of five to eight frames, and a screen that hides the table for two to sixteen frames while a swap is in motion. On five held-out worlds the belief patch points to the ball’s cup after 0.948 of the visible shuffles and on 0.890 of the last frames under the screen, 0.85 to 0.94 at every stretch length from two to sixteen frames; a record patch at matched parameters, whose context is a function of the input of the moment, points right on 0.545 of the visible shuffles and 0.541 under the screen; a nearest-neighbour tracker that extrapolates under the screen, the designed control, is the strongest arm on visible shuffles (0.953) and falls with the stretch, from 0.925 at two frames to 0.708 at sixteen, because a swap that completes under the screen ends on a slot and a constant velocity does not. Selection over the patch’s genome (its widths, repair iterations, damping, the store’s size and code, the rate) found one repair iteration in place of two at 106,549 multiply-adds per frame against 183,456, and that genome points right after 1.000 of the visible shuffles and on 0.992 of the last frames under the screen with no failed world, against 0.972 for the best of a random search at the same evaluations. The control wins on compute outright, about a thousand multiply-adds per frame; the store’s projection is four fifths of the patch’s cost [R114].

Habits distilled in imagination. A habit is a settling patch whose answer is its motor state, taught in the body model’s imagination without a teacher outside the brain. On the cart-pole, with the body a temporal patch learned from executed pushes, distillation from the imagination-fitted gains by the quadratic nudge on an opposing motor pair gives a habit that keeps the pole up through all 500 decisions of the shove run on every one of ten seeds, at a level error of 0.082 against its teacher; reward in imagination, the three-factor rule against imagined rollouts with the imagined cost as the punishment, drops the pole within 5 to 25 decisions on every seed after 200,000 imagined decisions and four world-judged rounds [R120]. On a two-link arm under gravity the same route reaches at 0.24 to 0.37 milliseconds and 1,300 to 9,700 multiply-accumulates per decision where the shooting planner it was distilled from costs 7.6 to 19.7 milliseconds and 983,040, the wall times taken on a loaded machine and the operation counts exact; a weight hung on the hand is caught by the task’s misses, the planner covers while the body learns the loaded arm from paths under the habit, and the habit is refitted in imagination [R121]. On Pong the route fails: a habit that reads the belief patch’s state reproduces the planner’s action on 0.936 of held-out beliefs (a linear readout 0.942) and loses every point, as does a second round taught on its own situations with the planner’s labels (0.918), because the decisions that return the ball are the rare exact presses that the belief’s play under either distilled policy never reproduces; label accuracy over a belief stream is a different quantity from the skill [R122].

A general recurrent F_θ need not arise from a scalar energy. Applying equilibrium detuning to such a candidate requires a compatible energy, stable teaching phases and a common fixed evidence and memory snapshot. Contraction of an inference map alone does not establish the required gradient identity.

2.7 Metacognition as recursive self-observation

Every patch above reads the world: its evidence is an observation, a body reading or a neighbor’s belief about the world across a port. Nothing in the contract prevents a patch from taking as its evidence the readback of other patches, their belief summaries, repair residuals and surprises, and returning its settled state to them as boundary conditions. We call such a patch a *steering patch*, and the number of readings between a patch and the world its *depth*.

Definition 2.1 (depth). A patch whose evidence consists of observations and of neighbors’ beliefs about the world has depth 0. A patch whose evidence includes the readback of other patches has depth one more than

the greatest depth among the patches it reads. The depth of a brain is the greatest depth of its patches. The depth of a task is the least depth of a brain that solves it at matched information and compute.

Depth is functional only under a condition, and the condition is a theorem. Let $x = \text{act}(h)$ be the published activity, let level 0 of a tower of readbacks publish C_0x , let level k publish C_k times the reading of level $k - 1$, and let every level re-enter the drive additively through a matrix D_k .

Theorem 2.2 (a linear tower has depth zero). *For every depth d ,*

$$Ax + \sum_{k < d} D_k C_k \cdots C_0 x = \left(A + \sum_{k < d} D_k C_k \cdots C_0 \right) x.$$

*The drive, the equilibria and every settling trajectory of the depth- d patch coincide with those of a depth-0 patch with the collapsed weights.*¹¹

The one-level case is the identity behind the readback control of Section 6 [R97]. A steering patch therefore contributes nothing unless it does one of three things: it publishes a summary that is nonlinear or lossy in what it reads; it acts on the equations of the patches below it, as a gain that multiplies their evidence, a gaze that changes which observation arrives, a goal that changes a clamped target or a budget that changes how long they repair; or it carries a cost of its own, so that its settled state is a choice and its residual and surprise are readback in turn.

The steering patch of Section 5 does the second. Its settled state is a weighing of two senses that multiplies each sense's encoded evidence before the repair map and the record reading see it, in every repair iteration of the moment, and the adjoint of the same equations returns the gradient into the gains, the seam through which the cortex's loss teaches the patch. Two free gains failed the condition: the patch drove both to the ceiling of their range in every phase and won by amplifying the evidence against the belief, a reparameterization the cortex could learn through its own repair map, and the phases were not separated; gains that sum to two removed the confound [R118].

The hypothesis. Metacognition is recursive self-observation: the settled state of a steering patch is the brain's reading of itself, and that reading steers the rest. Attention is a feature of it. A steering patch is bounded, so it cannot hold the whole equilibrium below it at once; it must select the aspects that matter for what it is computing, and that selection, of which evidence counts, where a sense samples, what a habit holds and how long a repair runs, is attention. Its own cost prefers states in which the surprise below it is consequential and falling, so that sensing and computation go to what is learnable and what matters; that preference is curiosity. Because the steering patch is a patch, its residual and surprise are readback in turn and a patch can read the patch that steers: the levels form the ladder of Section 5, and which readbacks exist is a question for the genome. The correspondences with the literature are these. Gain as the settled weight of evidence is precision weighting in predictive coding [14]. A large surprise below the steering patch that resets its state is the phasic locus coeruleus response to unexpected uncertainty [9, 53, 67], and the two directions of the selection correspond to the goal-directed and stimulus-driven systems of Corbetta and Shulman [11]. A patch that reads the steering patch has a confidence in a decision that is a reading of the process that made it [15], a model of where attention is and will be, the attention schema that a learning agent needed in order to control a spotlight of attention at all [61, 64], and an allocation of repair rounds and imagined branches by their expected value, which is metareasoning [33, 52]. A patch that reads a modeled observer, another agent or the brain's own earlier self, has the structure of a false-belief task [44, 62] and of dialogue. The linear read of Theorem B.5 is a different use of the word attention and establishes nothing about this hypothesis.

What a rung unlocks. A task requires depth d when brains of depth $d - 1$ fail it at matched information and compute and a brain of depth d passes. The ladder of Section 5 orders task families by the level conjectured to unlock them, and its first two rungs are earned there. With the existence of each readback a gene, the prediction is that selection adds a level exactly at the tasks the ladder places there and nowhere below. A shallow brain that solves a task placed higher, or a deeper brain that adds nothing on it at matched compute, falsifies that rung. The causal test for any level is the one of Section 6: at fixed observations, state and compute, remove or shuffle the level and measure the change in held-out prediction or return.

¹¹Lean: `readback_collapse`, `tower_eq_mulVec`, `drive_eq_collapsed`, `equilibrium_tower_iff`, `settle_tower_eq`.

3 Architecture and controlled evidence

The temporal experiments extend the patch interface to observed past state, private future state, action and preference ports. Their usefulness is measured by executed consequences and causal controls. Settling accuracy is reported separately from prediction, retention and control accuracy.

3.1 Temporal overlap carries credit only when the constraints carry time

In **TemporalPatchNet**, temporal patch t holds hidden state h_t and public readout y_t . Its transition/readout energy is

$$E_\theta = \frac{1}{2} \sum_t (\|h_t - A\phi(h_{t-1}) - Bu_t\|^2 + \|y_t - C\phi(h_t)\|^2). \quad (2)$$

Adjacent patches overlap on the preceding state; the initial cue u_t is absent after its first event. This residual family is established temporal predictive coding [35]. Detuning the readouts toward a complete observed sequence lets discrepancy propagate across temporal overlaps, so local parameter derivatives can receive credit from later observations.

Merely allocating future patches provides no choice. For deterministic transition and readback maps and a fixed initial state, every finite defect-free path equals the forward recurrence. Lean checks this without a linearity assumption¹². The statement concerns vanishing defects, not arbitrary stationary points. Useful planning therefore needs free action ports or additional constraints that make alternative consequences relevant [R83].

In the validated nonlinear sequence comparison, eight hidden coordinates and 80 learned parameters receive 1,536 updates over 16-, 32- and 64-event sequences. A distributed rotation initialization passes the 64/256-event continuation criterion in three development models, but fresh validation passes only one of three initializations on each of three-, four- and six-note motifs. A random orthogonal control passes one of nine. These are separately trained models, with every scheduled outcome retained. Each model receives 114,688 scalar teaching values in complete observed sequences, only during the detuned phases. Independent checks reconstruct derivatives, saved phase witnesses and autonomous outputs; they do not replay every training trajectory. The result supports finite temporal credit and some autonomous sequence acquisition, with neither reliable transfer nor online long-delay learning established [R83].

Coupled components share the same energy. Partitioning $h_t = (c_t, m_t)$ replaces A in Equation (2) by

$$A = \begin{pmatrix} A_{cc} & A_{cm} \\ A_{mc} & A_{mm} \end{pmatrix}; \quad A = M_A \odot W_A,$$

with analogous optional boolean masks for B, C . Both components settle jointly. For batch size N and transition defect e , the derivative for a legal block is $-N^{-1} \sum_{n,t} e_{nt}^a \phi(h_{n,t-1}^b)^T$, multiplied by its mask before the centered contrast and candidate admission. Concatenation is exactly the ordinary recurrence; component names supply no new capacity or specialization. Sequentially training detached components would omit this joint credit. The implemented solver retains dense blocks at total hidden width; zeros do not establish a computational saving [R96].

A bounded comparison gives both a two-group and a dense 64-coordinate model four supplied cyclic rotations of one recorded four-bar phrase, an initial cue, one silent primer and a periodic eight-position clock. A shared CPU cap stops both after 334 updates; all required phases converge and all 668 updates commit. Mean taught sample-identity accuracy is 29.69% for the partitioned model and 76.56% for the dense model, with 4,064 versus 7,104 trainable coefficients. Neither meets the complete recall criterion. Independent replay retains both capped outcomes. This single-seed, unequal-parameter comparison does not establish useful specialization [R96].

Initialization affects the available temporal signal. Masking the common orthogonal initializer lowers its spectral radius from 0.95 to 0.812. A separate, untrained diagnostic rescales every surviving recurrent weight by one common factor to restore radius 0.95. Fourth-bar cue-output Jacobian norm rises from 0.000295 to 0.0277, compared with 0.0354 for the dense initializer. This repairs an initialization mismatch, not the trained

¹²Lean: `defectFree_unique`, `eq_forward`, `readback_eq_forward`.

model’s failure. In general cue sensitivity obeys $S_t = A \text{diag}(\phi'(h_{t-1}))S_{t-1}$: a uniformly contractive path loses a one-time cue, whereas learning can change that path. Spectral radius alone bounds neither nonlinear memory duration nor the masked matrix’s operator norm. The diagnostic establishes larger finite cue sensitivity, not learned timescales or indefinite memory [R96].

3.2 The equilibrium chooses an action using an acquired body relation

The control body obeys $v_{t+1} = v_t + a_t + w_t$ and $p_{t+1} = p_t + v_{t+1}$; the actor reads consecutive positions, not hidden velocity. Inherited coordinates $h_t = (p_t, p_t - p_{t-1})$ expose position and displacement. Six parameters $W = [F \ G]$ learn the next state from real position readings and executed actions. For $E(z) = \|z - Wx\|^2/2$ with quadratic teaching, centered detuning is exactly ordinary squared-regression descent rescaled by $1/(1 - \beta^2)$. At $\beta = 0.1$, all three learned models achieve three-step held-out position MSE below 1.3×10^{-9} after 2,048 transitions. Matched no-history models fail. This is a deliberately transparent acquisition control, not a novel optimizer or learned coordinate system [R84].

With acquired F, G held fixed, future states and actions jointly minimize

$$E^g = \frac{q}{2} \sum_{t=0}^{H-1} \|h_{t+1} - Fh_t - Ga_t\|^2 + \frac{1}{2} \left\| h_H - \begin{pmatrix} g \\ 0 \end{pmatrix} \right\|^2 + \frac{r}{2} \sum_{t=0}^{H-1} a_t^2, \quad q = 1000, \quad r = 0.02. \quad (3)$$

The current observed boundary is fixed; g is a supplied retained preference. Adjacent three-coordinate patches exchange precision matrices and vectors. Only the first settled action is executed, after which the real position replaces the predicted readback and the actor replans. No demonstrated action sequence or external branch selector chooses the command. Soft dynamics constraints can bend imagined physics, so the evaluation measures actual executed cost and prediction defects rather than internal goal satisfaction alone.




Figure 1: The tested actor’s information flow. Observed records constrain the past; a copied estimate bounds private future repair. A preference changes the proposed action without rewriting the past estimate. Real consequences close the loop. The body relations are learned from separate actual experience and then fixed for this control evaluation.

Controller with acquired body	No disturbance	Hidden disturbance
Three-step feedback repair	0.009675	0.020846
One-step feedback repair	0.034174	0.055940
Initial plan, no later feedback	0.009675	1.234675
Three-step feedback, untrained body	16.38489	17.78535

Table 2: Mean executed cost $(p_3 - g)^2 + v_3^2 + 0.02 \sum_t a_t^2$ on the expanded control grid. Three independently acquired bodies are evaluated at $p_0 = \pm 0.225$, $v_0 = \pm 0.30$, $g = \pm 0.90$ and unannounced first-action disturbance $w_0 \in \{0, \pm 0.35\}$. Lower is better. The disturbance column pools both signs. The retained goal, deadline and physical coordinates are supplied [R84].

All three learned models pass the declared control criteria. Masking new observations reproduces the open-loop recovery failure. The independent audit checks 1,872 episodes across the original and expanded grids and 4,680 plans, including cold restoration, goal reversal and unchanged live state during private planning. Learned feedback actions do not saturate. Forecast error after the disturbed action is 3.01×10^{-10} ; this excludes the unpredictable first disturbance from the forecast-accuracy measure while preserving its full observed error. These controls establish that acquired relationships, future repair and real readback each contribute to action. They do not establish autonomous goal formation or protected memory during further learning [R84].

3.3 Continuous action repair through learned nonlinear dynamics

The same temporal energy supports action search without a separately supplied transition matrix. With parameters and initial boundary fixed, let e_t^+ and e_t^- be transition defects after positive and negative detuning toward a supplied goal. For batch size N , the centered input derivative is

$$g_t = -\frac{(e_t^+ - e_t^-)B}{2\beta N}, \quad L = \frac{1}{2} \text{mean}_{n,t,j} w_j (y_{ntj}^{\text{free}} - y_{ntj}^*)^2,$$

using row-vector defects and positive supplied output weights w_j . Temporal and output normalization enter through the teaching loss. Only declared control ports move; measured ports and the boundary remain fixed. Projected-gradient proposals obey the action bounds, and backtracking accepts one only when an ordinary target-free replay strictly decreases L and meets the Armijo condition. Under smooth stable-branch assumptions the contrast approaches the free-loss derivative as $\beta \rightarrow 0$; finite detuning and solver tolerance introduce error. The returned path is private, with separate reports for projected stationarity and predicted goal attainment. A goal-detuned future is never admitted as a measured outcome [R92].

A fully observed scalar body supplies actual transitions

$$p_{t+1} = 0.85p_t + 0.2 \tanh(2a_t) - 0.05p_t^3, \quad |a_t| \leq 1.$$

Its equation is used only by the environment. A sixteen-coordinate temporal learner receives an initial position, a presence flag and executed actions; subsequent position targets are real readings. Each of three fresh initializations receives 256 updates on batches of eight six-step paths, with 320 learned parameters. Held-out forward mean-square errors are 0.002102, 0.001183 and 0.001648. During control, weights stay fixed, only the first proposed action is executed, and each new measured position initializes another private plan. This fully observed encoding permits a cold hidden boundary; it does not test partially observed memory.

Controller with learned nonlinear body	No disturbance	Hidden disturbance
Six-step feedback proposal	0.025250	0.025520
One-step feedback proposal	0.025295	0.025565
Zero action	0.180744	0.194487
Random bounded action	0.276985	0.249866

Table 3: Actual eight-step mean squared goal error, averaged across three learned models. Initial position is ± 0.45 , goal is ± 0.35 , and an unannounced position displacement 0 or ± 0.25 occurs before the fifth action. Both feedback horizons use the same learned dynamics; lower is better [R92].

All three models meet the declared forward-prediction and executed-control criteria. Independent arithmetic reconstructs the data, transitions and scores; archived code separately replays all 768 learning updates and 576 plans across 144 episodes. All 288 six-step proposals reach the 32-step cap, despite improving predicted cost; they are not certified stationary optima. One-step control performs similarly on this constant-goal task. These results support learned consequence-guided action and disturbance correction, without establishing anticipation, curiosity, continued adaptation or a recursive self-model [R92]. On a two-link arm under gravity, with the body a 32-unit temporal patch learned from random torques, the gradient planner over that body did not reach its targets (a final hand distance of 0.2 to 1.8 metres across horizons, caps and rates) while random shooting over sixty-four candidate torque sequences in the same imagination reached them in 4 to 18 milliseconds; the body’s fault was the reconstruction of readings that sit far from zero, and centring and scaling the readings at the port cut the first imagined step’s hand error from 9 to 15 centimetres to 6 [R121].

A sampled-instrument control learns the rendered waveform’s interval RMS amplitude from sample identity and gain, then privately repairs continuous gains toward supplied acoustic goals. One of three scheduled models completes the bounded assessment; the other two have partial or unattempted outcomes, so there is no three-model success claim. With that fixed acquired model, a predetermined gain disturbance is repaired toward an unchanged reference recording: mean squared normalized RMS error falls from 0.09621 to 0.0001712, without training or disclosure of the target gains. Independent checks reproduce the private plan and executed sound. This is restoration under an externally supplied preference. Accurate control of an acoustic goal does not establish that the goal is musically desirable, that the actor learned the listener’s preference, or that it composes original phrases [R95].

3.4 Repair noisy interpretations while protecting measured records

An observation is a fixed record; an interpretation of it can be uncertain. The observer repairs latent position/displacement under actual position records and executed actions. Its Gaussian incoming message summarizes evicted history, with each measurement admitted once. The current mean and covariance are copied before the private goal/action solve. The observer has no goal argument, so a desired future cannot rewrite its factual estimate.

The tested model fixes acquired F, G , sensor variance $R = \sigma^2$, initial mean zero and covariance $100I$, and process covariance $Q = 0.04GG^\top + 10^{-6}I$. These are supplied model assumptions, not inferred reliability or a claim of optimality for the test disturbance distribution. Clean observations are hard position constraints, with a possibly singular posterior covariance. Under noisy readings, repairing the past improves pooled executed cost relative to raw position differences:

Sensor standard deviation	Raw differences	Repaired past
0.02	0.021238	0.021008
0.05	0.042065	0.036596

All 1,296 episodes are retained. At the smaller noise level, cost worsens in six of 27 model/noise/disturbance strata, by at most 0.92%; at the larger level every stratum improves. Independent all-history checks recover means and covariances to 3.91×10^{-13} and 1.22×10^{-15} and reject duplicate-evidence and goal-rewrites-past mutations. This is useful conditional state estimation, not immunity to arbitrary noise or false strong observations [R85].

3.5 A smaller sufficient boundary under a fixed model

For a fixed linear-Gaussian model, the Gaussian prefix is a sufficient boundary for later state inference. A one-record observer is therefore compared with a three-record observer on 27 fresh 64-transition streams. The current mean, covariance and proposed three-step actions agree to 9.05×10^{-13} , 3.39×10^{-14} and 5.83×10^{-12} . The observer’s persistent numeric payload, including a 32-byte model binding, falls from 272 to 208 bytes. Primary past factorizations fall from 5,265 to 1,782; additional prefix marginalizations rise from 1,701 to 1,755. Both kinds of work count. These exogenous streams test equivalence and cost, not another executed-control improvement [R85].

The compressed prefix binds to the values, shapes and dtypes of F, G, Q, R and the fixed prior. Changed models are rejected before the prefix can be silently reinterpreted. An independent sequential Gaussian oracle checks every prefix, with 378 additional dense full-history comparisons. Goal reversals preserve all past estimates and raw records; all 432 attempted operations after model/noise changes reject without modifying the prefix. This is conditional sufficiency at fixed model coefficients. It does not solve how to revise a learned model after discarding the observations that informed it [R85].

A scalar elimination theorem makes the information boundary exact. With fixed coefficients, $p > 0$, $q \geq 0$, and any later objective L that reads only y , define

$$E(x, y) = px^2 - 2bx + q(y - ax - u)^2 + L(y), \quad D = p + qa^2 > 0.$$

Completing the square gives

$$E(x, y) = D \left(x - \frac{b + qa(y - u)}{D} \right)^2 + p'y^2 - 2b'y + c' + L(y),$$

$$p' = \frac{pq}{D}, \quad b' = \frac{q(pu + ab)}{D}, \quad c' = \frac{qpu^2 + 2bqau - b^2}{D}.$$

The reduced expression is an attained lower bound in x . Thus two coefficients preserve later minimization over the retained boundary; the constant also preserves absolute energy. Admitting a new quadratic observation adds its precision and precision-weighted value to these coefficients. Six Lean theorems verify the identities and positivity assumptions¹³. The scalar proof and matrix experiment are separate evidence; neither is a theorem that an evolving nonlinear brain can compress all useful history this way [R86].

3.6 Retention and revision are stricter than equilibrium

A local response can be retained by constraining future parameter changes on its supporting activity. For a protected causal path, collect the presynaptic vectors for each shared map: $\tanh h_{t-1}$ for A , u_t for B , and $\tanh h_t$ for C . Let the orthonormal columns of Q_W span the selected vectors for map W . The closest update to a proposed detuning step D_W that leaves these local responses unchanged is

$$D_W^{\text{protected}} = D_W(I - Q_W Q_W^T), \quad D_W^{\text{protected}} Q_W = 0.$$

This is a quadratic parameter-overlap constraint, related to established orthogonal weight modification [68]. Preserving the transition maps on the old path preserves that path by induction; preserving its readback map preserves the emitted values. Three Lean lemmas state this conditional implication without axioms¹⁴. They assume exact map agreement and do not certify the floating-point SVD, relevance selection or nearby unprotected queries.

The bases compress local response constraints, with no raw phrase or target lookup during inference. The implementation applies ordinary centered detuning and the constrained update atomically. It binds the bases to the admitted parameter history and reports consumed ranks. Protection is selected by the caller and preserves the current learned response, including its existing error. Rank exhaustion can block further learning. Importance, selective fading and justified release of conflicting constraints remain separate questions.

3.7 Musical actions as a controlled example

A bounded real-audio task gives two streams the same eight heard bars and a single initial sign cue, followed by sixteen unobserved bars. The two intended futures are recorded phrases selected by position in pre-existing training splits from two supplied mixes. Future input is zero; targets enter only detuned learning phases. Descriptors use supplied whole-region normalization, so this is not a strict online-amplitude test. A 64-coordinate temporal network has 12,416 learned parameters. Across three initializations, all models meet the declared acquisition, cue-strength, temporal-variation, private-branch and cold-restoration checks. Mean continuation errors are 0.01423, 0.01319 and 0.01329. These are training-phrase realizations with supplied intention, not held-out-source generalization or autonomous goal formation [R88].

An additional third phrase tests further learning while preserving the first response. The protected path consumes ranks 23, 8 and 24 for A , B and C , or 28,328 numeric bytes including tolerance and parameter binding. The teaching cost uses per-output precision

$$w_j = \frac{\text{Var}_t(y_{tj}^{\text{new}}) + 0.0025}{\text{mean}_k[\text{Var}_t(y_{tk}^{\text{new}}) + 0.0025]},$$

computed from the observed new training phrase. These supplied weights add 512 checkpoint bytes; they are task weighting, not learned importance. The evaluation errors and variation gates are unweighted. After 1,536 updates on the new phrase, all three protected models pass the declared acquisition, retention, cue,

¹³Lean: `precision_pos`, `elimination_identity`, `reduced_attained`, `reduced_lower_bound`, `boundary_coefficients`, `admit_observation`.

¹⁴Lean: `protected_defectFree`, `protected_forward`, `protected_readback`.

variation, private-branch and cold-restoration checks. Maximum changes to the old response are 3.92×10^{-14} , 1.02×10^{-14} and 7.44×10^{-15} in independent recurrence checks. New-phrase errors are 0.006226, 0.004330 and 0.003606; variation ratios are 1.464, 1.279 and 1.139, within the fixed $[0.25, 4]$ interval. All three matched unprotected models learn the new phrase but change the old output by 0.992–1.173. Independent checks validate 36 saved phase endpoints and reject eight resealed mutations; they do not replay every intermediate update. This is conditional retention during additional phrase learning, not general lifelong memory [R89].

A separate sampler-grounded test learns actions that produce two recorded one-bar sounds. Sixteen supplied half-beat crops contain the two bars at their native 175 beats per minute. Twenty literal action coordinates specify sample choice, presence, signed offset, gain and duration. A temporal network with 64 hidden coordinates and 5,504 parameters receives a shared activation and an intention sign only at the first of eight events; later input is zero. Source-derived action sequences supply teaching targets, and known-action playback checks the sampler before training.

After 1,536 updates, all three development seeds meet the declared recall, cue, timing, restoration and private-imagination checks. All 48 sample choices across the two intentions and three seeds are correct; reversing the cue selects the other taught sequence. Independent recurrence, saved-phase and exact audio replay checks bind the learned action path to its sound. The listening files use the prescribed first seed and repeat one generated bar four times. Timing within each crop is inherited from the recording. This establishes conditional sample-sequence recall and sound realization; it does not establish original composition, auditory-feedback learning, fresh-seed generalization or retention after learning another task. A learned instrument relation, useful higher-level readback and coherent long-form development are separate requirements [R90].

A four-bar repertoire uses 32 recorded slices, a two-coordinate initial cue and a 64-coordinate temporal learner with 6,592 parameters, jointly trained for 1,536 updates. All three models meet the declared recall criteria, with 96 of 96 taught sample choices correct. Three untaught cue combinations produce mixed slice orders; all are rendered from the prescribed first seed without ranking. Independent checks reproduce their PCM and MP3 files exactly. The task has no absolute mean-square-error gate. It tests joint repertoire acquisition and cue interpolation, not sequential retention or deliberate composition [R93].

Response protection can leave usable but poorly conditioned directions. In a three-program sampled-action task, protection consumes ranks 14, 2 and 16 for A , B and C , or 15,432 numeric bytes, leaving 4,160 parameter directions available. The fully constrained input map fixes the cold first activity q . If Q spans protected readout activity and $r = (I - QQ^T)q$, every correction with $\Delta CQ = 0$ and $(C + \Delta C)q = y^*$ satisfies

$$\|\Delta C\|_F \geq \frac{\|y^* - Cq\|_2}{\|r\|_2}.$$

Residual activity norms of 0.0245–0.0322 require readout changes of norm 21.9–28.8 for exact first-output correction. This bound concerns conditioning, not rank exhaustion or whole-program impossibility [R91].

A readout metric addresses this numerical constraint without teaching from protected targets. Let X contain free presynaptic activities, $S = X^T X / (NT)$, $P = I - QQ^T$ and g_C be the same centered detuning contrast. With supplied damping $\mu > 0$, the update

$$\Delta C = -\eta g_C P (PSP + \mu I)^{-1} P$$

uniquely minimizes $\langle g_C, D \rangle + \text{tr}[D(S + \mu I)D^T] / (2\eta)$ subject to $DQ = 0$. The conditioning matrix uses free activity and protected bases; teaching enters through the same contrast; A and B retain ordinary projected detuning. This is a patch-level dense solve, not an independent per-synapse operation. The metric uses S , whereas the fixed-feature half-mean-loss Hessian is S/O for O output coordinates. Damping and scaling are declared choices, not learned importance [R94].

The matched test uses the same acquired two-bar models and 1,536 updates on a supplied third sample program in both arms, with $\beta = 0.001$, $\eta = 1$ and $\mu = 10^{-4}$ for the metric. All three metric models preserve both old responses within 4.74×10^{-14} and fit the complete new action path below 3.48×10^{-23} mean squared error. All three ordinary projected models preserve the old responses but have new errors 0.008546, 0.011507 and 0.011048. The new-path thresholds are 0.001 and one quarter of its initial error. Thus the metric passes all three joint retention/acquisition criteria; ordinary projected detuning passes none. Every model chooses all sixteen old and eight new sample identities correctly, illustrating why category accuracy alone is insufficient. The prescribed first seed preserves both old rendered WAV files byte for byte. Independent checks replay every

one of the 9,216 updates, all 27,648 phase records and twelve PCM files, with separate constrained-optimum and finite-difference witnesses [R94].

This result has a narrow numerical scope. The metric adds one 64×64 solve per update and no persistent optimizer state; readout changes have norm 32.88–42.20, compared with 1.62–1.68 for ordinary updates. Off-path noise sensitivity is unmeasured. Positive curvature at both detuned endpoints does not establish a common stable branch: at a fixed checkpoint, $\beta = 0.1$ produces an ascent direction despite positive endpoint curvature. At the final near-zero losses, relative gradient errors are dominated by solver tolerance. The matched test has no parameter backtracking; one seed was used in development and two extend that fixed choice. Protection and the new program are supplied. The result establishes conditional acquisition with retention, not automatic importance, selective forgetting, unlimited capacity or original composition [R94].

4 Language against a next-word learner

A transformer learns language as next-token prediction: one forward pass over the whole window per token, knowledge in weights moved by a gradient over a corpus, a new fact reaching it through the window or through further gradient steps, and a cost per token that grows with the window. Table 4 states where the patch net of this paper differs; the differences are the same for language and for reasoning, and the evidence below concerns language.

	next-word learner	patch net
state	none between tokens; the window is re-read at every step	a bounded settled state carried between events (working memory), and a record store for what does not fit; the window is the focus
memory	in the weights, by gradient over the corpus	one write per observation into records (Theorem B.1); the slow weights take what recurs
learning signal	cross-entropy of the next token	local repair: the delta rule for records and the free/nudged contrast for the slow weights; for the record patch the contrast equals the adjoint gradient
production	one token per forward pass	the design is a settled sentence: all word slots consistent with the message before the first word is read out; the measured speaker reads word by word from the store
understanding	implicit in the window’s activations	a settled state consistent with the evidence and the constraints, readable at ports: a message exists before its words
reasoning	longer windows and more passes	constraint satisfaction over the persistent state under a focus; the free phase over the whole store with nothing clamped is the offline form of the same operation

Table 4: A next-word learner and the patch net, operation by operation. The production, understanding and reasoning rows describe the design; the measured speaker reads word by word from the store, and no reasoning task is measured in this paper.

Controlled evidence on a finite grammar of 1,276 messages and 6,173 sentences supports parts of this design and not others. Next-word prediction itself is not what separates the two: taught by cross-entropy, the patch ties with a GRU and a transformer. What the patch has and they lack are the other two routes, one write per observation and the night.

Depth is needed for syntax. One record patch, whose gate sees only the present word, chooses the verb’s number correctly for 1.00, 1.00 and 0.03 of sentences with zero, one and two nouns of the other number between subject and verb; two stacked patches, a GRU and a two-layer transformer choose it correctly for 1.00 throughout. The same single patch learns an arbitrary function of the ordered last two symbols from its slow weights alone (0.9999), because the input-dependent gate multiplies the carried context [R102].

Records generalize without a gradient. With the slow weights left random, one pass of delta-rule writes over 10,672 taught pairs, about a minute, gives a sentence inside the grammar for 0.815 of the (subgenre, era,

form) combinations never taught and 0.885 of the other held-out messages; the grammatical member of a minimal pair (a/an, is/are, deal/deals, both/all) is preferred in 0.9975 and 1.00 of them. A second pass gives 0.85 and 0.915 [R103].

With enough data the three learners tie. Trained by cross-entropy for fifty epochs, a two-layer patch, a GRU and a transformer all reach 1.00 on the same held-out sets; the patch is the largest (9.68M against 3.90M and 4.70M parameters) and the slowest (2,920 against 748 and 887 seconds) [R105]. At a twelfth of the pairs and equal updates the patch has the lowest held-out loss of the three at every width (0.32 against 0.34 and 0.36 nats per word at width 64); at a quarter it alone is at 1.00 valid at every width; at half and at all of the pairs the three learners lie within 0.03 nats of one another [R104].

The store is for one-shot learning. Twenty-four sentences the slow weights never saw are each recited exactly after one write, and none with the store off. After all twenty-four are written, 0.79 of them are recited exactly and the 2,765 read sentences fall from 1.000 to 0.973: a shared store disturbs what it holds, and a conversation needs its own store [R106].

4.1 Acquisition in two phases

The tested record patch separates acquisition into two phases. By day, observations update the records (Theorem B.1) while the slow weights are held fixed. By night, every supplied cue is completed once by the store; the slow weights learn those fixed dreams without writing new records, and at dawn the dreams are written back relative to the learned readout. The replay targets therefore come from the patch’s retained state. This is one implementation inspired by the fast and slow learning distinction of complementary-learning-systems theory [30, 34]. It does not require all slow learning to occur during sleep, and the comparison concerns this acquisition protocol rather than a universal distinction between patches and transformers.

In this protocol, the day’s targets survive only through the store, so repeated presentations use its completions. Learning their shared structure can generalize beyond individual cues, but also trains the store’s errors. Retaining observed targets would permit a different replay control. The protocol does not establish that generated dreams are necessary for generalization.

Measured on the same grammar, in Wagner’s design [59]: at bedtime the store alone gave 0.815 of the never-taught combinations a valid sentence and the slow weights, untrained, none (one knob at a time, 64 active cells in place of 32 lift the store’s share to 0.900, a write rate of 0.5 lowers it to 0.675 and a store of 16,384 cells to 0.655 [R112]). After one night of 8,000 dreams from the day’s own cues (0.9095 of them inside the grammar) and 1,200 slow updates with the corpus closed, the slow weights alone gave 1.00, and so did the patch with its store; the awake control that re-read the corpus twice more with the teacher present reached 0.905 [R108].

Sleep also carried the store’s errors. On the held-out messages whose completions were wrong at bedtime the night changed nothing (0.885 to 0.89), where the teacher with the corpus for the same 1,200 updates reached 0.98. Dreams from recombined cues, which land on combinations never taught, gave 0.83 and 0.875: the store’s guesses at unmet cues are taught as facts, whereas dreams from the cues the store met let the smooth weights extend the pattern. Sleep supplies no independent evidence that a completion is correct. A record is the residual against the slow readout at the time it was written, so the dreams are fixed at bedtime and the store is re-referenced at dawn; dreaming again after each update teaches the weights their own drift [R108]. Two corrections without the corpus were tried: keeping only the dreams whose every word won by a margin of 0.3, which kept 13 of 8,000, since the share inside the grammar was the same among the unsure dreams as overall (0.909); and re-reading the unsure cues from the teacher at dawn, which re-read 7,987 of them and reached the teacher’s 0.98. The store’s own certainty does not sort its right dreams from its wrong ones [R113].

The same regime was run on three of the published examples with the library’s sleep unchanged. The Connect Four value patch, with its school closed after one day of 67,399 writes, went from chance to 0.755, 0.768 and 0.772 on the held-out sign after one night at 4,096, 16,384 and 65,536 cells, against 0.785 for the school from the same 263 updates on the outcomes themselves; at every matched update count the school was one to three points ahead, the dawn readout followed what the store read at bedtime (0.733 to 0.789), and a second night on the same store added nothing, since its dreams were what the weights said; on the full

school of 4.4 million positions taken as 67 days and nights the sleep-trained readout reached 0.838 against the school's 0.896 at the same 52,000 updates, so the gap grows with the school [R109]. The composer, with its corpus closed after one day, reached in one night of 3,833 updates the held-out next-event accuracy that three epochs with the corpus reached in 11,499: 0.826 to 0.852 for the slow parameters alone against 0.828 to 0.850, and 0.831 to 0.864 with the records writing online against 0.843 to 0.872; the second and third nights moved nothing [R110]. In the soft-body world a sleep gene that keeps the day's windows and dreams them by night was selected out in every run, from 0.37 of the founders to 0.01 to 0.03 of the living after 20,000 ticks, with a late-life model error of 0.056 to 0.058 against 0.035 to 0.049 for the beings that learn by day: an awake being has the world at every tick, and a sleeper trades its actions for second-hand presentations of what its store holds [R111]. These comparisons establish that the value of sleep depends on target quality, access to observed outcomes and the cost of displaced interaction. They do not establish a universal separation between learning while awake and learning during sleep.

Three things are not established. The settled sentence is a design: a speaker built as one equilibrium net with twenty winner-take-most word slots coupled through an association layer, taught by clamping the message and nudging every slot at once, reached a slot accuracy of 0.19 after one epoch over 2,668 pairs with settles that did not converge in sixty steps, below the 0.5 of always saying nothing, and was stopped [R107]. No reasoning task is measured; the reasoning row of Table 4 is the design. Nothing here compares with a transformer at scale: the grammar is finite, generalization means unseen combinations inside it as judged by an acceptor, and every learner reaches the ceiling with enough data.

5 The metacognition ladder

The end of the program is a brain in a body that functions among people in the real world: it looks, acts and speaks, holds a plan, asks when unsure and says what it is doing; past that, brains that read their own development and one another. The route is a ladder of steering patches. Each rung is a function that the brain below it lacks and a task on which that lack is measured, with a control and a falsifier declared before the run. Table 8 states the rungs. Rungs 1 and 2 are earned by the demos of this section; every rung above them is a hypothesis, and a rung is earned only by a task the brain below fails at matched information and compute.

The demo rule. Every step of the program is accepted by a demo. Its page carries the whole brain in the viewer, the world, an instrument strip (the repair residual, the surprise, the update, the gains sent down, the moments per decision), a three-way switch (the step on; off, the brain below it; the hand-designed control) and the person in the world: the mouse as the laser dot, the hand that drags the voice, the hand that shuffles the cups. A step is accepted when the demo works with the switch on and fails with it off; the control keeps both of its outcomes. Every parameter that looks designed is a gene, selected against the hand-set value, which stays the control, with a random search at the same number of evaluations as the second control. Table 5 states the four acceptances under this rule, two at the foot of the ladder and one on each of its first two rungs.

Selection over the genome. In all four demos selection over a dict genome, with the hand-set genome as the lineage's first member and a random search over the same space at the same number of evaluations, beat both on held-out worlds: the cat's governor 0.853 against 0.782 and 0.642 on the fitness, the ventriloquist's steering patch 0.038 against 0.041 and 0.056 on the error, the shell game's belief patch 0.992 against 0.890 and 0.972 under the screen, the ball's ports 0.169 against 0.230 and 0.219 in the dark [R116, R118, R114, R115]. Selection also pruned what did not pay: the cat's lineage switched learning off, the ventriloquist's dropped the residual probes and the repair residual from the readback and kept the two surprises alone, the ball's dropped the port into the ear, and the shell game's dropped the second repair iteration. In the room below, where a fixed reflex holds the task on its own, the random search's winner, which never imagines and never learns, beat the lineage under both fitnesses; both outcomes are kept.

The foot of the ladder. The measured brains read the world only. The *C. elegans* connectome runs as one masked temporal patch net, its browser brain and body held to the library and to a reference body, and a supplied worm mechanism is probed under lesions [R98, R37]. Soft bodies with two record patches live under selection in a world of conserved mass, and the share of sleepers under a sleep gene is followed over

demo (rung)	the step on	off: the brain below	the hand-designed control	compute: on / off / control
the shell game (0): the belief patch's private continuation an eye and an ear (0): record patches joined by ports	pointed cup 0.948 visible, 0.890 under the screen; evolved genome 1.000 and 0.992 eye error in the dark 0.169 with the evolved port from the ear	record patch at matched parameters: 0.545 and 0.541 0.310 with that port cut; 0.280 with no port	nearest-neighbour tracker: 0.953 and 0.825, falling to 0.708 at sixteen frames hand-designed ports both ways: 0.230	multiply-adds per frame 183,456 (106,549 evolved) / 113,436 / 1,008 [R114] parameters 3,717 / 3,692 / 3,634 at one budget; an update with the hand ports 0.138 s against 0.063 s without [R115]
the dozing cat (1): the governor patch	catch rate 0.961, awake 0.238 of the time	always awake 0.962, awake 1.000; never wakes 0.249	threshold rule: 0.931 hand-set, 0.909 evolved	patch moments per decision 15.3 / 71.3 and 1.0 / 26.2 (17.5 evolved) [R116]
the ventriloquist (2): the steering patch	error 0.041 (0.038 evolved) over the three phases	fixed gains 0.081; 0.086 widened to matched compute	inverse-variance rule: 0.062	multiply-accumulates per decision 63,470 (55,657 evolved) / 54,400 (63,584 widened) / 54,406 [R118]

Table 5: The four acceptances under the demo rule: what the demo does with the step on, what the brain below it does with the step off, what the hand-designed control does, and the compute of each, on held-out worlds. Every number is a mean over five worlds (three for the eye and the ear).

twenty thousand ticks [R100, R111]. A mouse connectome finds its way through a labyrinth against a shuffled connectome and a network [R24]. A value patch schooled by a perfect player plays Connect Four with a supplied search that imagines moves and reads the patch where it stops looking, with the search alone as its control and the day and night regime on the same patch [R101, R109]. A records controller flies the 1942 shooter against a convolutional policy at matched frames, and the 1943 imitation player practices on 1942 [R66, R69]. The composer recalls and realizes taught bars [R90, R99]; the critic that selects among its genomes reads rendered audio and never the brain. The language learner of Section 4 reads a corpus. In none of these does a patch read another patch. The shell game and the eye and the ear of Table 5 stand here too: the belief patch that carries the ball under the screen and the ports that keep the eye's belief alive in the dark read the world only [R114, R115].

5.1 Rung 1, earned: the dozing cat

A cat dozes on a sill on one moment of its belief patch per decision (Figure 2). The belief patch reads a six-by-six retina, the fovea's centroid, optic flow and mass, and its own paw, 43 readings, and predicts the change of the dot's centroid and mass and of the paw: 6,853 parameters, 96,488 multiply-accumulates per moment, pretrained on 120 streams of random pushes with spans of moments hidden so that the transition carries the belief without evidence. A dot appears on a random schedule as a walk, a curve or a jitter and is caught when the paw rests within reach for three decisions; at decision 2,400 of 4,000 its law changes for good to gravity. Three modes: the habit (the paw drifts home, one moment per decision), imagination (ten candidate pushes held for six decisions in the belief's private imagination, sixty moments, the best first push executed) and learning (the belief observes its executed window of 96 decisions, the update kept when the window's loss fell by a tenth, the habit refitted in imagination). The readback port publishes every decision the last surprise over its baseline, log-compressed, a slow average of the same, the repair residual over its routine median, the last mode and a constant. The governor patch, the rung, is a settling patch of the same neuron model as every other brain here: a port of seven units, a cortex of four and a motor of three whose most active unit is the mode; every synapse and bias is a gene, 77 in all, and its settle costs 23 synapses times about twenty steps, 0.005 moments per decision. The hand-designed control is a threshold rule with twelve genes: a spike above k baselines recruits imagination for a budget of decisions, and persistence above a share of the last 96 triggers learning. The brain below the rung comes in two forms, never waking and always awake. The fitness is the catch rate minus 0.004 per moment per decision minus 0.5 per unit of the mean surprise

capped at one, every channel at weight one. Table 6 gives five held-out lives per arm.




Figure 2: The dozing cat’s page: the sill with the cat and the laser dot; the whole brain in the viewer, with the governor patch’s readback port, cortex and motor drawn as their own regions beside the senses, the evidence, the belief, the records, the prediction and the habit; the three-way switch under it; the mode, the awake share, the catches and the moments per decision. The page runs the brain in the browser with the library’s arithmetic [R117].

arm	catch rate	awake share	awake, no dot	moments/decision	learn calls (kept)	fitness
governor patch, evolved (on)	0.961	0.238	0.031	15.3	0.0 (0.0)	0.853
governor patch, hand-set	0.977	0.365	0.222	37.3	13.8 (3.6)	0.782
threshold rule, hand-set (control)	0.931	0.271	0.144	26.2	6.6 (2.6)	0.783
threshold rule, evolved	0.909	0.269	0.134	17.5	0.2 (0.2)	0.795
always awake (off)	0.962	1.000	1.000	71.3	6.4 (3.4)	0.631
never wakes (off)	0.249	0.000	0.000	1.0	0.0 (0.0)	0.167
random search’s best, patch genome	0.857	0.205		42.1		0.642
random search’s best, threshold genome	0.952	0.283		29.4		0.788

Table 6: The dozing cat on five held-out lives of 4,000 decisions with the law changing at decision 2,400: the catch rate (caught over caught or missed), the share of decisions awake, the same when no dot is on the sill, the patch moments per decision, the learn calls per life with the kept updates in parentheses, and the fitness [R116].

With the switch on, the evolved governor patch catches 0.961 of its dots awake 0.238 of the time (0.80 with a dot on the sill, 0.03 without), waking within a decision of a dot’s appearance, at 15.3 moments per

decision. Off, the brain below fails both ways: never waking catches 0.249 and misses twelve dots a life; always awake catches the same, 0.962, at 71.3 moments per decision, 4.7 times the cost, awake all the time. The threshold rule catches 0.931 hand-set and 0.909 evolved at 26.2 and 17.5 moments. The patch’s lineage climbed from 0.824 to 0.906 on its training life in six generations of twelve, and its winner scores 0.853 on the held-out lives against 0.782 for the hand-set wiring and 0.642 for the best of 72 random draws; it moved 65 of its 77 genes, and what it did is readable in them: the fast surprise drives the imagine unit harder, the imagine unit’s self-excitation is gone, the cortex the hand-set wiring left empty carries the slow surprise and the current mode, and the learn unit’s drive is halved under a bias, so the evolved governor never learns. The threshold lineage found nothing beyond its first generation and beats its hand-set values by 0.012. Both lineages switched learning off, and the receipt says why: under gravity the dot is slower than the paw, every governed arm catches every dot after the change with or without learning, so a kept update buys prediction alone, worth 0.0015 of fitness through the surprise term, against 0.056 of fitness for the observe passes and the refit, about fourteen moments per decision. A task on which the changed law costs catches, or a surprise price thirty times higher, would price learning in. The falsifier of the rung, that the settling governor does no better than hand-set thresholds at matched compute, does not fire: at the hand-set genomes the two tie on the fitness (0.782 against 0.783, the patch catching more at more compute), and after selection over each genome’s own space at an equal budget of 72 lives the patch wins on every part. The page runs the evolved genome, and the same brain runs in the browser with the library’s arithmetic as the published example at <https://floatingpragma.io/cadence-examples/dozing-cat/>, held to the library by a parity test whose worst relative difference is 4.8×10^{-12} on the committed recording [R116, R117].

The room, the governor’s compute and two negatives. The first page of the rung keeps a room at 21 degrees through a heat pump: one belief patch reads the temperature, its rate and the light; doors kick the rate, lights flicker, and at decision 2,400 a window opens for good. The same three modes under a threshold governor with 22 genes cost 21.6 moments per decision, against 56.0 imagining every decision and 154.8 learning on a schedule of every 96 decisions, at equal task cost, since the fitted habit holds the room on its own; the governor makes no learn call over the 48 doors and lights of a life before the change on any seed, detects the window 73 decisions after it opens and an aged heater after 98, and one kept update brings the surprise down where the never-learning arm’s stays up. Under a price on compute alone selection switches the governor off: the random search’s winner, which never imagines and never learns, scores -0.0098 held out against -0.0151 for the lineage and -0.0222 for the hand-set genome. Under a price on surprise as well the lineage keeps the whole governor, learns on every held-out seed and imagines after events at 6.9 moments per decision, and the deaf reflex, at a surprise term of 0.222 against 0.084, wins at the price used; the break-even is a surprise price of about 0.034 per decision. A fitness that read the surprise the governor itself computes, weighted by the genome’s own channel weights, was gamed: the random search’s winner set the rate’s weight to 0.018 and the light’s to 0.053 and reported a surprise term of 0.024 while never learning, so a fitness that prices surprise reads it with every channel at weight one [R119]. The same governor on Pong, with a belief patch of 3.74 million parameters and the rules swapped at decision 2,000 of 6,000, sees nothing: the planner holds one direction on 80 percent of decisions and the paddle sits at a wall, where it does not move under any action, so the belief’s expectation of its own paddle is met in both regimes and the retina surprise (0.047 before, 0.038 after) and the residual stay where they were while the score drops to the floor for every arm. A policy that saturates its actuator hides the change of an action’s meaning from the surprise; only a brain that moves can be surprised by how it moves. One kept update at decision 940, valid by the window’s one-step loss, wrecked the planner’s open-loop imagination, which nothing in the validity check measured [R119].

5.2 Rung 2, earned: the ventriloquist

A dummy moves at a steady speed in a box and bounces off the walls (Figure 3). An eye sees it as a six-by-six retina and two coordinates, all scaled by the light; an ear hears the voice as a four-channel interaural cue. One belief patch reads both senses through a structured port with a block per sense (38 inputs to 16 encoded units for the eye, 4 to 8 for the ear), a belief of 32 units, two repair iterations and a store of 256 cells, and predicts the dummy’s next position: 6,810 parameters, 54,400 multiply-accumulates per decision. Three phases follow a hidden schedule the brain never sees: in the light the eye and the ear agree; when the ventriloquist speaks the voice is displaced from the dummy by a hidden offset and the eye is right; in the dark the retina is blank and the ear is right. The steering patch is a second belief patch of the same rule over a readback port of five

channels, the repair residual of each sense alone, the surprise of each sense (what it read against what the cortex’s predicted position implies it should have read, in units of the sense’s persistence error) and the repair residual of the moment before: 438 parameters, 1,326 multiply-accumulates. Its two outputs pass through a weighing, $2\text{softmax}(\sigma \tanh y)$ with a span σ , so the gains sum to two and the patch can only move weight from one sense to the other; the gains multiply each block’s encoded evidence inside every repair iteration and the store’s code, and the patch learns within the life from the cortex’s loss through that seam. The controls: fixed gains at one, the brain below the rung; the same widened until its compute matches the steering arm’s; the designed inverse-variance rule, each gain the inverse of a running mean of its surprise; and two ablations of the trained steering brain at test time, the gains frozen at one and the gains swapped between the senses. Learning is 120 epochs on 48 episodes of 64 frames per world; the test is 32 held-out episodes with the brain frozen. Table 7 gives five worlds; Figure 4 the gains and the readback by phase.




Figure 3: The ventriloquist’s page in the light phase: the box with the dummy, the voice and the crosshair of the selected brain’s next place; the whole brain in the viewer, the cortex’s retina, ear, evidence, belief, records and prediction beside the steering patch’s readback port, evidence, belief and the two gain units whose synapses return to the evidence; the gain sliders, moving on their own; the three-way switch [R118].

With the switch on, the steering patch weighs the eye at 1.79 against 0.21 when the ventriloquist speaks and the ear at 1.78 against 0.22 in the dark, and its error is 0.041 overall, 0.049 in the ventriloquist phase and 0.036 in the dark, at the level of the phase oracle (0.040), a ridge told the phase that reads the right sense, and below stay-put (0.059) in every phase. Off, the same cortex with fixed gains is fooled where the ladder says it should be: 0.081 overall, 0.124 when the ear pulls the belief toward the displaced voice and 0.077 when the blank retina pulls it toward what a blank retina encodes, above stay-put in both; widened to the steering arm’s compute it stays fooled at 0.086. The inverse-variance rule moves the gains the right way and reaches 0.062, behind the steering patch in every phase. The ablations within a life go the way the rung requires: the trained steering brain with its gains frozen at one degrades to 0.117 (0.220 in the dark) and with its gains swapped to 0.147, worse than the brain trained with fixed gains, because the cortex that grew up under the steering patch never learned to gate its senses itself. The readback shows what the patch works from: the surprise of the displaced ear (2.75 against 0.51 in the light) and of the blank eye (4.39 against 0.24); the residual probes barely move and the repair residual not at all (0.24 to 0.25). Selection over the steering patch’s genome (its

arm	overall	light	ventril.	dark	gains eye/ear: ventril.	dark	MACs/decision
steering patch, hand-set (on)	0.041	0.039	0.049	0.036	1.79 / 0.21	0.22 / 1.78	63,470
steering patch, evolved	0.038	0.037	0.041	0.036	1.83 / 0.17	0.24 / 1.76	55,657
inverse-variance rule (control)	0.062	0.045	0.093	0.058	1.41 / 0.59	0.43 / 1.57	54,406
fixed gains (off)	0.081	0.056	0.124	0.077	1.00 / 1.00	1.00 / 1.00	54,400
fixed gains, widened	0.086	0.057	0.136	0.082	1.00 / 1.00	1.00 / 1.00	63,584
steering patch, cut at test	0.117	0.042	0.136	0.220	1.00 / 1.00	1.00 / 1.00	63,470
steering patch, shuffled at test	0.147	0.052	0.165	0.285	0.74 / 1.26	1.32 / 0.68	63,470
random search's best	0.056	0.043	0.095	0.039	1.44 / 0.56	0.14 / 1.86	62,652
stay-put	0.059	0.059	0.059	0.059			
phase oracle	0.040	0.043	0.043	0.033			

Table 7: The ventriloquist on the held-out episodes of five worlds: the distance between the predicted and the true next position, overall and per hidden phase, the gains the brain gives the eye and the ear when the voice is displaced and in the dark, and the multiply-accumulates per decision. Stay-put predicts no motion; the phase oracle is a ridge on the last two readings of the right sense [R118].




Figure 4: The ventriloquist's gains by phase on the held-out episodes of five worlds: the eye's and the ear's gain under the steering patch, the inverse-variance rule and fixed gains (hollow marks per world, filled marks the mean, the dashed line the neutral weighing), and the surprise readback the steering patch reads, in units of each sense's persistence error [R118].

widths, iterations, damping, span, rate and which readback channels it reads) kept the two surprises alone and dropped the probes and the residual, which drops the probes' 7,744 multiply-accumulates per decision with them: the evolved brain reaches 0.038 at 55,657 multiply-accumulates, 2.3 percent above the fixed brain's, where the hand-set genome's costs 63,470, 17 percent above; the best of 24 random draws, which reads the probes and no surprise, scored above every member of the lineage on its one training life and 0.056 on the test worlds, with its gains pinned to the ear in every phase on one of them. On the page the same happens under the mouse: the voice dragged away moves the weighing to 1.88 against 0.12, with the fixed brain missing by 0.144 to the steering patch's 0.037, and the lights out to 0.21 against 1.79, with 0.133 against 0.044 [R118]. The falsifier of the rung, fixed gains matching held-out prediction at matched compute, does not hold. The gain multiplies the evidence inside the settle, which is what Theorem 2.2 requires of a level that is more than a reparameterization; a linear mixing of two outputs after the settle would collapse.

Table 8: The metacognition ladder. Each rung names what the steering patch reads and returns, what it is conjectured to unlock, and the test on which the brain below it is predicted to fail at matched information and compute. Rungs 0 to 2 are measured; every rung above 2 is a hypothesis.

rung	reads; returns	unlocks	test
0	the world; predictions and actions	tracking, habits, one-shot recall, grammar completion, control of a learned body	the measured brains of this paper: the connectome worm, the soft-body world, the mouse labyrinth, Connect Four, the 1942 pilot, the composer, the language learner, the shell game, an eye and an ear
1	its own surprise; its mode, habit, imagine or learn	learning only when needed, imagination on demand	earned: the dozing cat, a brain that runs on habit in a quiet room and wakes to imagine when a moving cue appears, measured as decisions awake against catches (Table 6); a drummer’s tempo change caught by surprise is untested
2	the residuals of several senses; their evidence gains	trusting the right sense at the right time	earned: two senses report one dummy and one is displaced in phases; the eye is weighed when they disagree in the light and the ear in the dark (<i>the ventriloquist</i> , Table 7)
3	the residual map; the gaze of a sense under a sensing cost	frames larger than the patch repairs; selective sensing	Pong and Seaquest seen through a window the brain moves (<i>the periscope</i>); a beam swept over a dark sea with ships arriving from random bearings (<i>the lighthouse keeper</i>)
4	a large surprise below seizes it; repetition habituates	orienting, the interrupt, habituation	an attended conversation and a rare rewarded cue in the ignored one (<i>the cocktail party</i>); a ticking clock habituates while an alarm that predicts fire keeps capturing
5	its own cost prefers falling, consequential surprise; where the senses go	learning what matters without a schedule	a screen of static beside a slow cartoon (<i>the noisy television</i>); toys of different learnability and the order in which they are left (<i>the toy box</i>)
6	the residual and the stakes; repair rounds and imagination budget under a deadline	cheap routine, deliberate hard cases	Connect Four with a time bank per game (<i>the chess clock</i>); a crossing game in which hesitation kills
7	the steering patch; a confidence, an opt-out, a request for help	knowing what it does not know	the labyrinth mouse that stops to look when its place is uncertain, scored by metacognitive sensitivity; a quiz with a costly lifeline; a wager on its own decision
8	the steering patch; a prediction of its own gaze and gain, said in words	the attention schema; grounded language	a brain that narrates its own Pong game, graded against its logged gaze, with a narrator that sees the frame and not the steering patch as control (<i>the sports commentator</i>); the composer’s diary of its intentions
9	imagination itself; which futures to imagine, how far, when to stop	strategy, a plan before the first move	a key, a door, a vault and a guard on patrol (<i>the heist</i>); boxes pushed into corners and lost (<i>Sokoban</i>); Connect Four searched by its own uncertainty against uniform search
10	the goal-holder; a goal for the goal	long-form creation, multi-stage tasks, the desirability of goals	a four-minute arc that every bar serves, rated blind against the bar-level composer (<i>the DJ set</i>); a composition that every stroke serves (<i>the mural</i>); a dungeon of rooms and keys
11	a modeled observer’s evidence; its belief	false belief; another’s attention	a seeker that models where the hider thinks it will look (<i>hide and seek</i>); a brain that misdirects another brain’s gaze (<i>the magician</i>)
12	another brain’s sentences; a shared belief	world models by words	two brains with different views of one world describe and ask, scored by joint prediction against silent brains (<i>the two lookouts</i>)
13	the student’s surprise; the next lesson	teaching, learning from teaching	a competent brain teaches a fresh one and says the lesson, against fixed and random curricula (<i>the driving instructor</i>)
14	all of the above on a real body	the real world	an obstacle course crossed through a moving gaze in simulation, then fetch with a person on a real quadruped
15	all of the above among people	the end of the program	a humanoid that makes a coffee and says what it is doing (<i>the kitchen helper</i>); a shelf assembled with a friend; a tour given to a visitor
16	its own development; other brains’ steering patches; deeper nesting	past people	a brain that proposes the next experiment on itself; a hundred brains sharing one sense (<i>the parliament</i>); a hidden-information game played at a depth of belief people do not sustain

What the ladder predicts. Three predictions follow from the hypothesis and are the ladder’s falsifiers. First, on each rung’s task the brain of the rung below fails at matched information and compute and the brain with the rung passes; a shallow brain that passes retires the rung, and a deeper brain that adds nothing retires it too. This holds on rungs 1 and 2: never waking and always awake fail the sill and the governor patch passes; fixed gains fail the ventriloquist and the steering patch passes, with the hand-designed rule between them. Second, with the existence of each readback a gene, selection on a rung’s task adds a steering patch at that rung and nowhere below; a random search at the same evaluations is the control. This is tested in part: with the channels the steering patch reads as genes, selection kept the surprise readback and dropped the probes and the residual, and in the cat’s wiring it moved the readback into a cortex the hand-set wiring left empty; a search in which the existence of the level itself is a gene is work in progress. Third, with a developmental clock a gene, a brain whose deeper readbacks open late in a life, in the order of the ladder, outperforms a brain whose readbacks are open from birth; Section 6 gives the developmental order this prediction is drawn from, and it is untested. If every rung above the second collapses under these tests, the hypothesis is false and the route to the end of the program is a different one.

6 Biological abstractions and functional self-readback

Animal and human brains motivate distributed recurrent interpretation, different memory timescales, learned action consequences and specialization through unequal experience. These are functional design constraints, not a claim that Cadence reproduces the brain’s anatomy or dynamics. Persistent-activity theories motivate carried context [19, 60]; complementary-learning-systems theories motivate separating fast experience from slower shared structure [30, 34]; predictive-coding models motivate local discrepancy signals [7, 48]. Each abstraction must earn its inclusion by retaining information or improving a measured decision. The two-phase acquisition of Section 4.1 is one such test: the store retains the day, and the night’s replay of its own completions gives the slow weights what they do not produce alone.

Functional self-readback requires a causal role. Feeding $y = C\phi(h)$ back into $h' = A\phi(h) + Dy$ is equivalent to $h' = (A + DC)\phi(h)$; it does not create additional evidence or memory. An explicit matrix control verifies this reparameterization [R97]. The control experiment instead compares predicted consequences with new physical readings and demonstrates that masking those readings impairs recovery. The steering patches of Section 5 are a higher-level view of this kind: they read surprises, residuals and the mode and return a mode and gains. Their usefulness requires a matched causal intervention: holding observations, state budget and compute fixed, remove or shuffle that view and measure the behavioral change. Section 5 applies it: the steering patch cut at test raises the ventriloquist’s error from 0.041 to 0.117 and shuffled to 0.147, and the governor patch removed leaves the cat asleep at a catch rate of 0.249 or awake at 4.7 times the cost [R118, R116]. Recursively stacking interfaces alone proves no such function, and Theorem 2.2 states the linear case exactly.

Depth in development. The order in which these functions appear in children is the order of the rungs of Section 5. Alerting and orienting are present in infancy, disengaging fixation becomes reliable in the third and fourth month, and executive attention emerges at the end of the first year and matures through childhood [46, 51]. Infants of twelve to eighteen months monitor their own errors and persist longer after correct choices [20], and at twenty months they ask for help when they do not know [21]. Mirror self-recognition appears between eighteen and twenty-four months [2, 49]; explicit false-belief understanding appears in the fourth and fifth year [62], with an implicit form debated from fifteen months [44]. The anatomy follows the same sequence: synaptic density peaks in auditory cortex at three months and in prefrontal cortex after fifteen months, with pruning there into adolescence [25]; higher-order association cortex matures after the sensory cortex it integrates [18]; the default network is adult-like by about two years [17]; and functional networks go from local to distributed [13]. For the hypothesis this suggests that the readbacks of deeper levels open late in a life, in the order of the ladder, which a genome can encode as a schedule and selection can test against a brain whose readbacks are open from birth. Higher-order theories of awareness place the reading of one’s own states in prefrontal cortex [32], and Cleeremans et al. [10] propose that the brain learns to redescribe its own activity to itself, a formulation close to the ladder, whose fixed point is the strange loop of Hofstadter [23]. This paper makes no consciousness claim; the ladder is a hypothesis about function, tested by ablation.

Specialization must also be useful. Exchangeable patches can receive unequal experience or break initial symmetry, but unequal weights or a named region do not establish division of labor. A learned division should improve behavior under causal lesion and transfer tests while preserving unrelated records during further learning. The fixed-model boundary and supplied-address retention results state concrete cases where information can be separated; discovering the required separation in music is a different problem.

7 Limits and the generalization test

The control results cover acquired linear dynamics with inherited coordinates and a fully observed scalar nonlinear body, each with supplied goals and acquired dynamics frozen during execution. Longer planning helps the linear task but supplies no established anticipation benefit on the scalar nonlinear task. Gaussian compression is exact only relative to its fixed linear model. The real-audio results use complete teaching sequences, three initializations and selected training phrases; held-out musical generalization is unestablished. Exact response protection is conditional on selected local constraints and usable plastic directions; the sampled-action comparison shows that a supplied readout metric can address poor conditioning despite unconsumed rank. Its large coefficients and fixed damping require further robustness tests. These are not general lifelong-memory guarantees.

The language results hold on a finite grammar judged by an acceptor, against a GRU and a transformer of comparable size at equal updates, not at scale. The settled-sentence speaker failed its first calibration, the tested sleep protocol did not reliably repair the store’s wrong completions, and reasoning is a design row with no measurement. Rungs 1 and 2 are earned on one world each. The cat’s changed law costs no catches, so learning was priced out and the evolved governor never learns; the ventriloquist’s senses are declared functions of the position, which is what makes the surprise readback computable, its task is one-step prediction, and the hand-set steering patch costs 17 percent more compute per decision than the fixed brain (2.3 percent for the evolved genome). The controls keep their wins: the nearest-neighbour tracker at a thousandth of the belief patch’s compute, the deaf reflex in the room at the prices used, the wide patch on the ear’s bounces, and on Pong both the distilled habit and the governor’s surprise. Every wall-clock number of the demos was taken on a loaded machine; the moment and operation counts are exact. Every rung of Table 8 above the second is a hypothesis, and Theorem 2.2 states what a level must do to be more than a reparameterization. The world-model proposal in Section 2.6 requires separate evidence for nonlinear composition, calibrated action consequences, imitation beyond the demonstration distribution and fully internal continuation. Neither the fixed-point theorems nor the finite grammar establish these capabilities or an advantage over attention-based models.

Generalization requires useful behavior across new circumstances, retained experience and changed body consequences. For composition, the decisive test is joint: a brief intention must survive a shared heard prefix; changed intention must change the continuation; altered instrument consequences must change the next plan; imagined branches must leave witnessed state untouched; and learning new material must preserve and, when appropriate, revise old material. Context length, autonomous variation and low one-step error are separate measurements. Short intention-conditioned phrase acquisition and explicit response protection establish parts of this standard. The complete learned instrument, self-readback and long-form composition loop is unestablished. Training on the available jungle corpus and creating original whole tracks is an intended application, not a result demonstrated by these one-bar controls. The same intervention standard must hold in other sensory and action domains; adding modalities supplies no universal scaling guarantee.

Total cost includes all free and detuned phases, failed numerical attempts, matrix factorizations, communication, model acquisition, inference and any retained evidence. Local parameter derivatives do not imply low wall-clock cost. A conventional regression, filter, recurrence or controller is a necessary comparison when it solves the same stated problem. Learned developmental wiring, importance selection and robust transfer across tasks require evidence beyond the fixed small worlds measured here.

8 Conclusion

Cadence seeks a small set of reusable operations for continuing intelligence: repair interpretations, learn relationships through detuning, and compare proposed consequences with actual readback. The proposed world

model represents the actor within its surroundings and tracks familiar situations through a changing sequence of consistent states. Recursive nonlinear composition, learning from demonstrations and interventions, and grounded consolidation are candidate means of making that trajectory useful and inexpensive. New evidence can revise the current belief without changing the underlying law; private imagination explores proposed continuations without becoming witnessed evidence. Controlled bodies and musical actions establish narrower instances of repair and action. On a finite grammar, a record store takes a corpus in one pass without a gradient, and one night of its own completions transfers useful structure into slow weights. Explicit memory constraints and sufficient boundaries state what can be preserved; a matched readout-metric comparison demonstrates further acquisition while retaining two constrained responses. Under the demo rule the first two rungs of the ladder are earned: a governor patch that reads the brain’s own surprise catches the dots on a sill awake a quarter of the time, and a steering patch that reads a cortex’s readback weighs the right sense in every phase, each against the brain below it and a hand-designed control, and in every acceptance demo selection over the genome beat the hand-set genome and a random search. Scalable world modeling, automatic importance, any rung of the ladder above its second, learned specialization, robust transfer and original whole-track composition are not established. The ladder orders the tasks between these results and a brain in a body among people, and each rung is earned by a task the brain below fails. Progress requires those joint capabilities at a declared total cost, beyond numerical equilibrium alone.

References

- [1] David H. Ackley, Geoffrey E. Hinton, and Terrence J. Sejnowski. A learning algorithm for Boltzmann machines. *Cognitive Science*, 9(1):147–169, 1985. doi: 10.1207/s15516709cog0901_7.
- [2] Beulah Amsterdam. Mirror self-image reactions before age two. *Developmental Psychobiology*, 5(4):297–305, 1972. doi: 10.1002/dev.420050403.
- [3] Jimmy Ba, Geoffrey E. Hinton, Volodymyr Mnih, Joel Z. Leibo, and Catalin Ionescu. Using fast weights to attend to the recent past. In *Advances in Neural Information Processing Systems*, volume 29, 2016. URL <https://arxiv.org/abs/1610.06258>.
- [4] Shaojie Bai, J. Zico Kolter, and Vladlen Koltun. Deep equilibrium models. In *Advances in Neural Information Processing Systems*, volume 32, 2019. URL <https://arxiv.org/abs/1909.01377>.
- [5] Bowen Baker, Ilge Akkaya, Peter Zhokhov, Joost Huizinga, Jie Tang, Adrien Ecoffet, Brandon Houghton, Raul Sampedro, and Jeff Clune. Video PreTraining (VPT): Learning to act by watching unlabeled online videos. *arXiv preprint arXiv:2206.11795*, 2022. URL <https://arxiv.org/abs/2206.11795>.
- [6] Stefan Banach. Sur les opérations dans les ensembles abstraits et leur application aux équations intégrales. *Fundamenta Mathematicae*, 3:133–181, 1922.
- [7] Andre M. Bastos, W. Martin Usrey, Rick A. Adams, George R. Mangun, Pascal Fries, and Karl J. Friston. Canonical microcircuits for predictive coding. *Neuron*, 76(4):695–711, 2012. doi: 10.1016/j.neuron.2012.10.038.
- [8] Josh Bongard, Victor Zykov, and Hod Lipson. Resilient machines through continuous self-modeling. *Science*, 314(5802):1118–1121, 2006. doi: 10.1126/science.1133687.
- [9] Sebastien Bouret and Susan J. Sara. Network reset: a simplified overarching theory of locus coeruleus noradrenaline function. *Trends in Neurosciences*, 28(11):574–582, 2005. doi: 10.1016/j.tins.2005.09.002.
- [10] Axel Cleeremans, Dalila Achoui, Arnaud Beauny, Lars Keuninckx, Jean-Remy Martin, Santiago Muñoz-Moldes, Laurène Vuillaume, and Adélaïde de Heering. Learning to be conscious. *Trends in Cognitive Sciences*, 24(2):112–123, 2020. doi: 10.1016/j.tics.2019.11.011.
- [11] Maurizio Corbetta and Gordon L. Shulman. Control of goal-directed and stimulus-driven attention in the brain. *Nature Reviews Neuroscience*, 3(3):201–215, 2002. doi: 10.1038/nrn755.
- [12] Chris Eliasmith. *How to Build a Brain: A Neural Architecture for Biological Cognition*. Oxford University Press, 2013. ISBN 9780199794546. doi: 10.1093/acprof:oso/9780199794546.001.0001. URL <https://academic.oup.com/book/6263>.
- [13] Damien A. Fair, Alexander L. Cohen, Jonathan D. Power, Nico U. F. Dosenbach, Jessica A. Church, Francis M. Miezin, Bradley L. Schlaggar, and Steven E. Petersen. Functional brain networks develop from a “local to distributed” organization. *PLoS Computational Biology*, 5(5):e1000381, 2009. doi: 10.1371/journal.pcbi.1000381.
- [14] Harriet Feldman and Karl J. Friston. Attention, uncertainty, and free-energy. *Frontiers in Human Neuroscience*, 4:215, 2010. doi: 10.3389/fnhum.2010.00215.
- [15] Stephen M. Fleming and Hakwan C. Lau. How to measure metacognition. *Frontiers in Human Neuroscience*, 8:443, 2014. doi: 10.3389/fnhum.2014.00443.
- [16] Nicolas Frémaux and Wulfram Gerstner. Neuromodulated spike-timing-dependent plasticity, and theory of three-factor learning rules. *Frontiers in Neural Circuits*, 9:85, 2016. doi: 10.3389/fncir.2015.00085.
- [17] Wei Gao, Hongtu Zhu, Kelly S. Giovanello, J. Keith Smith, Dinggang Shen, John H. Gilmore, and Weili Lin. Evidence on the emergence of the brain’s default network from 2-week-old to 2-year-old healthy pediatric subjects. *Proceedings of the National Academy of Sciences*, 106(16):6790–6795, 2009. doi: 10.1073/pnas.0811221106.
- [18] Nitin Gogtay, Jay N. Giedd, Leslie Lusk, Kiralee M. Hayashi, Deanna Greenstein, A. Catherine Vaituzis, Tom F. Nugent,

- David H. Herman, Liv S. Clasen, Arthur W. Toga, Judith L. Rapoport, and Paul M. Thompson. Dynamic mapping of human cortical development during childhood through early adulthood. *Proceedings of the National Academy of Sciences*, 101(21):8174–8179, 2004. doi: 10.1073/pnas.0402680101.
- [19] Patricia S. Goldman-Rakic. Cellular basis of working memory. *Neuron*, 14(3):477–485, 1995.
- [20] Louise Goupil and Sid Kouider. Behavioral and neural indices of metacognitive sensitivity in preverbal infants. *Current Biology*, 26(22):3038–3045, 2016. doi: 10.1016/j.cub.2016.09.004.
- [21] Louise Goupil, Margaux Romand-Monnier, and Sid Kouider. Infants ask for help when they know they don’t know. *Proceedings of the National Academy of Sciences*, 113(13):3492–3496, 2016. doi: 10.1073/pnas.1515129113.
- [22] Danijar Hafner, Jurgis Pasukonis, Jimmy Ba, and Timothy Lillicrap. Mastering diverse domains through world models. *arXiv preprint arXiv:2301.04104*, 2023. URL <https://arxiv.org/abs/2301.04104>.
- [23] Douglas R. Hofstadter. *I Am a Strange Loop*. Basic Books, New York, 2007.
- [24] J. J. Hopfield. Neurons with graded response have collective computational properties like those of two-state neurons. *Proceedings of the National Academy of Sciences*, 81(10):3088–3092, 1984. doi: 10.1073/pnas.81.10.3088.
- [25] Peter R. Huttenlocher and Arun S. Dabholkar. Regional differences in synaptogenesis in human cerebral cortex. *Journal of Comparative Neurology*, 387(2):167–178, 1997. doi: 10.1002/(SICI)1096-9861(19971020)387:2<167::AID-CNE1>3.0.CO;2-Z.
- [26] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. Transformers are RNNs: fast autoregressive transformers with linear attention. In *International Conference on Machine Learning*, 2020. URL <https://arxiv.org/abs/2006.16236>.
- [27] Teuvo Kohonen. Correlation matrix memories. *IEEE Transactions on Computers*, C-21(4):353–359, 1972. doi: 10.1109/TC.1972.5008975.
- [28] Dmitry Krotov. A new frontier for Hopfield networks. *Nature Reviews Physics*, 5:366–367, 2023. doi: 10.1038/s42254-023-00595-y.
- [29] Dmitry Krotov and John J. Hopfield. Dense associative memory for pattern recognition. In *Advances in Neural Information Processing Systems*, volume 29, 2016. URL <https://arxiv.org/abs/1606.01164>.
- [30] Dharshan Kumaran, Demis Hassabis, and James L. McClelland. What learning systems do intelligent agents need? complementary learning systems theory updated. *Trends in Cognitive Sciences*, 20(7):512–534, 2016. doi: 10.1016/j.tics.2016.05.004.
- [31] Axel Laborieux, Maxence Ernoult, Benjamin Scellier, Yoshua Bengio, Julie Grollier, and Damien Querlioz. Scaling equilibrium propagation to deep ConvNets by drastically reducing its gradient estimator bias. *Frontiers in Neuroscience*, 15: 633674, 2021. doi: 10.3389/fnins.2021.633674.
- [32] Hakwan Lau and David Rosenthal. Empirical support for higher-order theories of conscious awareness. *Trends in Cognitive Sciences*, 15(8):365–373, 2011. doi: 10.1016/j.tics.2011.05.009.
- [33] Falk Lieder and Thomas L. Griffiths. Strategy selection as rational metareasoning. *Psychological Review*, 124(6):762–794, 2017. doi: 10.1037/rev0000075.
- [34] James L. McClelland, Bruce L. McNaughton, and Randall C. O’Reilly. Why there are complementary learning systems in the hippocampus and neocortex. *Psychological Review*, 102(3):419–457, 1995. doi: 10.1037/0033-295X.102.3.419.
- [35] Beren Millidge, Mufeng Tang, Mahyar Osanlouy, Nicol S. Harper, and Rafal Bogacz. Predictive coding networks for temporal prediction. *PLOS Computational Biology*, 20(4):e1011183, 2024. doi: 10.1371/journal.pcbi.1011183.
- [36] Javier R. Movellan. Contrastive Hebbian learning in the continuous Hopfield model. *Connectionist Models: Proceedings of the 1990 Summer School*, pages 10–17, 1991. doi: 10.1016/B978-1-4832-1448-1.50007-X.
- [37] Bernhard Mueller. Cadence: machine learning by settling nets with owner-local learning, 2026. URL <https://github.com/muellerberndt/cadence>. Python library, MIT license; measured implementation versions and source hashes are identified in the accompanying evidence package.
- [38] Bernhard Mueller. Cadence memory: representation boundaries and local-repair diagnostics, 2026. Executable experiments, source-bound receipts, independent verifiers and Lean 4 proofs in the accompanying evidence package.
- [39] Bernhard Mueller. Verified projection-event calculus in Lean 4: Bundled arbitrary-partition pinching, Lüders retractions, and CHSH interoperability, 2026. URL <https://floatingpragma.io/oph/papers/machine-checked-finite-event-algebras/>.
- [40] Bernhard Mueller. Observer patch nets: local repair, retained state and the console evidence, 2026. URL <https://github.com/muellerberndt/reverse-engineering-reality>. Observer patch holography corpus.
- [41] Bernhard Mueller. Thinking as patch-net fixed-point search, 2026. URL <https://github.com/muellerberndt/reverse-engineering-reality>. Observer patch holography corpus.
- [42] Bernhard Mueller, Jinwook Kim, David Matscheko, and Jonathan Hill. Observation-determined normal forms: Stability, obstructions, and refinement in constraint and rewrite systems, 2026. URL <https://floatingpragma.io/oph/papers/observation-determined-normal-forms/>. Observer patch holography corpus, machine-checked in Lean 4 with Mathlib.
- [43] Bernhard Mueller, Kai Xue, Jinwook Kim, Arnav Anirudha Kale, David Matscheko, and Jonathan Hill. Reality as a consensus protocol: The fixed-point computation that implements physics, 2026. URL <https://floatingpragma.io/oph/papers/reality-as-consensus-protocol/>.
- [44] Kristine H. Onishi and Renée Baillargeon. Do 15-month-old infants understand false beliefs? *Science*, 308(5719):255–258, 2005. doi: 10.1126/science.1107621.
- [45] Randall C. O’Reilly. Biologically plausible error-driven learning using local activation differences: the generalized recirculation algorithm. *Neural Computation*, 8(5):895–938, 1996. doi: 10.1162/neco.1996.8.5.895.
- [46] Michael I. Posner and Mary K. Rothbart. Research on attention networks as a model for the integration of psychological science. *Annual Review of Psychology*, 58:1–23, 2007. doi: 10.1146/annurev.psych.58.110405.085516.
- [47] Hubert Ramsauer, Bernhard Schäfl, Johannes Lehner, Philipp Seidl, Michael Widrich, Thomas Adler, Lukas Gruber, Markus

- Holzleitner, Milena Pavlović, Geir Kjetil Sandve, Victor Greiff, David Kreil, Michael Kopp, Günter Klambauer, Johannes Brandstetter, and Sepp Hochreiter. Hopfield networks is all you need. In *International Conference on Learning Representations*, 2021. URL <https://arxiv.org/abs/2008.02217>.
- [48] Rajesh P. N. Rao and Dana H. Ballard. Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects. *Nature Neuroscience*, 2:79–87, 1999. doi: 10.1038/4580.
- [49] Philippe Rochat. Five levels of self-awareness as they unfold early in life. *Consciousness and Cognition*, 12(4):717–731, 2003. doi: 10.1016/S1053-8100(03)00081-3.
- [50] Stephane Ross, Geoffrey Gordon, and Drew Bagnell. A reduction of imitation learning and structured prediction to no-regret online learning. In *Proceedings of the Fourteenth International Conference on Artificial Intelligence and Statistics*, volume 15 of *Proceedings of Machine Learning Research*, pages 627–635. PMLR, 2011. URL <https://proceedings.mlr.press/v15/ross11a.html>.
- [51] M. Rosario Rueda, Jin Fan, Bruce D. McCandliss, Jessica D. Halparin, Dana B. Gruber, Lisha Pappert Lercari, and Michael I. Posner. Development of attentional networks in childhood. *Neuropsychologia*, 42(8):1029–1040, 2004. doi: 10.1016/j.neuropsychologia.2003.12.012.
- [52] Stuart Russell and Eric Wefald. Principles of metareasoning. *Artificial Intelligence*, 49(1–3):361–395, 1991. doi: 10.1016/004-3702(91)90015-C.
- [53] Susan J. Sara and Sebastien Bouret. Orienting and reorienting: the locus coeruleus mediates cognition through arousal. *Neuron*, 76(1):130–141, 2012. doi: 10.1016/j.neuron.2012.09.011.
- [54] Benjamin Scellier and Yoshua Bengio. Equilibrium propagation: bridging the gap between energy-based models and back-propagation. *Frontiers in Computational Neuroscience*, 11:24, 2017. doi: 10.3389/fncom.2017.00024.
- [55] Imanol Schlag, Kazuki Irie, and Jürgen Schmidhuber. Linear transformers are secretly fast weight programmers. In *International Conference on Machine Learning*, 2021. URL <https://arxiv.org/abs/2102.11174>.
- [56] Julian Schrittwieser, Ioannis Antonoglou, Thomas Hubert, Karen Simonyan, Laurent Sifre, Simon Schmitt, Arthur Guez, Edward Lockhart, Demis Hassabis, Thore Graepel, Timothy Lillicrap, and David Silver. Mastering Atari, Go, chess and shogi by planning with a learned model. *arXiv preprint arXiv:1911.08265*, 2019. URL <https://arxiv.org/abs/1911.08265>.
- [57] Wolfram Schultz, Peter Dayan, and P. Read Montague. A neural substrate of prediction and reward. *Science*, 275(5306):1593–1599, 1997. doi: 10.1126/science.275.5306.1593.
- [58] Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2 edition, 2018.
- [59] Ullrich Wagner, Steffen Gais, Hilde Haider, Rolf Verleger, and Jan Born. Sleep inspires insight. *Nature*, 427:352–355, 2004. doi: 10.1038/nature02223.
- [60] Xiao-Jing Wang. Synaptic reverberation underlying mnemonic persistent activity. *Trends in Neurosciences*, 24(8):455–463, 2001.
- [61] Taylor W. Webb and Michael S. A. Graziano. The attention schema theory: a mechanistic account of subjective awareness. *Frontiers in Psychology*, 6:500, 2015. doi: 10.3389/fpsyg.2015.00500.
- [62] Henry M. Wellman, David Cross, and Julianne Watson. Meta-analysis of theory-of-mind development: the truth about false belief. *Child Development*, 72(3):655–684, 2001. doi: 10.1111/1467-8624.00304.
- [63] Bernard Widrow and Marcian E. Hoff. Adaptive switching circuits. *IRE WESCON Convention Record*, 4:96–104, 1960.
- [64] Andrew I. Wilterson and Michael S. A. Graziano. The attention schema theory in a neural network agent: Controlling visuospatial attention using a descriptive model of attention. *Proceedings of the National Academy of Sciences*, 118(33):e2102421118, 2021. doi: 10.1073/pnas.2102421118.
- [65] Xiaohui Xie and H. Sebastian Seung. Equivalence of backpropagation and contrastive Hebbian learning in a layered network. *Neural Computation*, 15(2):441–454, 2003. doi: 10.1162/089976603762552988.
- [66] Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. In *Advances in Neural Information Processing Systems*, volume 37, 2024. URL <https://arxiv.org/abs/2406.06484>.
- [67] Angela J. Yu and Peter Dayan. Uncertainty, neuromodulation, and attention. *Neuron*, 46(4):681–692, 2005. doi: 10.1016/j.neuron.2005.04.026.
- [68] Guanxiong Zeng, Yang Chen, Bo Cui, and Shan Yu. Continual learning of context-dependent processing in neural networks. *Nature Machine Intelligence*, 1:364–372, 2019. doi: 10.1038/s42256-019-0080-x.

A Certified recurrence and local credit

These statements concern the declared recurrent subclass; a prototype must meet their hypotheses before using their guarantees.

A.1 Neurons, synapses, settling

A net has n neurons. Neuron i holds a potential $v_i \in \mathbb{R}$ and publishes an activation $s_i = \text{act}(v_i)$. Synapses carry effective weights W_{ij} from neuron j to neuron i ; in the library the effective weight is the product of a gain, the number of contacts, a learned efficacy and a presynaptic factor, and the paper writes the product as

W_{ij} . A stimulus c_i is an additive drive, with the bias absorbed into it. One settling step is

$$v_i \leftarrow (1 - dt) v_i + dt \left(\sum_j W_{ij} \text{act}(v_j) + c_i \right), \quad 0 < dt \leq 1. \quad (4)$$

The activation is a rectified, re-based logistic. With slope λ , threshold θ , leak $\ell \in [0, 1]$ and $\sigma(x) = 1/(1 + e^{-x})$,

$$\rho = \sigma(-\lambda\theta), \quad r(v) = \sigma(\lambda(v - \theta)) - \rho, \quad \text{act}(v) = \frac{\max(r(v), 0)}{1 - \rho} + \ell \frac{\min(r(v), 0)}{\rho}. \quad (5)$$

Rest publishes exactly zero, the activation rises to one, and a neuron below rest publishes a small negative value so that the learning rule can move it. The learning model uses $\lambda = 1$, $\theta = 0$, $\ell = 1/10$. A settling run stops when the largest activation movement in a step falls below a tolerance or when a step cap is reached, and the library reports the residual of equation (4) at the final state. Adaptation, a slow variable subtracted from the drive, produces rhythmic dynamics; the theorems below are stated without it.

A declared neural graph can include sensory, internal and motor coordinates. Joint settlement can complete a partial reading when its learned constraints support that completion; region names alone assign no function.

Intuition: each neuron checks its own account. Imagine an arm whose visual input has moved while its motor and body predictions describe the preceding position. Each neuron compares its held potential with the drive its neighbors and senses supply. Equation (4) moves the potential partway toward that drive; the published activation changes what neighboring neurons receive, and they make the same comparison. The resulting computation is a feedback process with a measurable discrepancy at every step. In the observer-patch description, each bounded patch owns state, communicates through declared ports, reads back its consequences and retains records. A settling certificate and the public evidence bundle make that self-reading process checkable. The equilibrium reconciles the equations the model implements; whether those equations describe the arm well is a separate empirical question.

A.2 From observable events to a global normal form

OPH distinguishes internal descriptions from shared readings [41–43]. A visual patch may locate an object in camera coordinates and a motor patch in body coordinates. Agreement concerns the position after translation to a shared interface. A change of internal coordinates and corresponding interface map is a gauge change when it preserves every declared public reading. Gauge-invariant repair corrects disagreement in those readings.

For local states x_i and readout maps $\pi_{i,e}$ into a common interface on an overlap $e = \{i, j\}$, agreement means

$$\pi_{i,e}(x_i) = \pi_{j,e}(x_j).$$

A repair preserves the observations held fixed. OPH’s local-repair criterion requires every allowed fixed exterior to admit a consistent interior. A *normal form* has no repair left to make. For a terminating discrete relation on the observable quotient, local confluence gives one normal form from each starting state. Completeness, meaning that normal states coincide with consistent states, makes it a global solution. Local compatibility alone need not imply global consistency. For different interiors with the same protected observations to yield the same answer, the compatible consistent states must also form one equivalence class [42, 43].

What counts as an observable event? For a finite classical record, events are sets of possible readings, closed under union, intersection and complement. OPH’s verified projection-event calculus supplies conditional expectations that discard internal distinctions while preserving the probabilities of the declared partition events [39]. Cadence uses classical activation ports and writable records; the quantum projection and conditioning machinery is mathematical background, without a claimed implementation in this neural model.

The neural specialization. Hold the weights fixed, protect $B(c, v) = c$, and define consistency by $v = W \text{act}(v) + c$. Under Cadence’s contraction condition, every initial state at a fixed input approaches the same equilibrium. The residual and stimulus bounds quantify accuracy and sensitivity. This realizes an observation-determined equilibrium: local equations determine one global answer with a checkable error bound. Numerical convergence replaces finite exact termination; the synchronous proof does not establish arbitrary asynchronous or transactional repair guarantees. Learning determines whether the answer is useful.

A.3 The certificate: how much work a decision needs

Write $F_c(v)$ for the right-hand side of (4) at stimulus c , $\|\cdot\|_\infty$ for the sup norm on $\mathbb{R}^n$, $\rho_W = \max_i \sum_j |W_{ij}|$ for the largest absolute incoming weight sum (the row mass), and L for a Lipschitz constant of act .

Theorem A.1 (settling certificate). *If $0 < dt \leq 1$ and $L\rho_W < 1$, then F_c is a contraction of $(\mathbb{R}^n, \|\cdot\|_\infty)$ with constant $q = 1 - dt(1 - L\rho_W) < 1$.*¹⁵

Corollary A.2 (unique equilibrium). *Under the hypotheses of Theorem A.1 the net has exactly one equilibrium v_c^* , and every fixed point of F_c equals it; with zero stimulus and $\text{act}(0) = 0$ that equilibrium is rest.*¹⁶

Theorem A.3 (a priori and a posteriori certificates). *Under the hypotheses of Theorem A.1, for every start v and every k ,*

$$\|F_c^k(v) - v_c^*\|_\infty \leq \frac{q^k}{1-q} \|v - F_c(v)\|_\infty, \quad \|F_c^k(v) - v_c^*\|_\infty \leq \frac{\|F_c^k(v) - F_c^{k+1}(v)\|_\infty}{1-q}.$$

17

A residual below ϵ at the returned state certifies a potential error of at most $\epsilon/(1-q)$, and an activation error of at most L times that. This is an error bound relative to the model's equilibrium, not a certificate of prediction accuracy. Banach's theorem is the whole proof [6]; the activation-specific condition is given in Theorem A.6. A numerical control checks the bound on 57,600 cases [R8].

Theorem A.4 (the equilibrium is Lipschitz in the stimulus, and warm starts pay). *Let F_c and $F_{c'}$ satisfy Theorem A.1 with the same constant q and let $\delta = \|c - c'\|_\infty$. Then*

$$\|v_c^* - v_{c'}^*\|_\infty \leq \frac{dt \delta}{1-q}, \quad \|F_{c'}^k(v_c^*) - v_{c'}^*\|_\infty \leq \frac{dt \delta q^k}{1-q} \quad \text{for every } k.$$

18

This is the theorem behind adaptive work. Starting from the old equilibrium, the number of steps that reaches accuracy ϵ after an input change δ is at most the nonnegative ceiling of $\log(dt \delta / ((1-q)\epsilon)) / \log(1/q)$: logarithmic in the size of the change, zero when the change is zero. The bound is sufficient rather than a prediction of semantic difficulty, and checking a returned state has its own cost.

Intuition: reuse the answer, then bound the repair. Contraction says that two possible internal states move closer under the same input: after one step their largest difference is at most q times its previous value. With $q = 1/2$, every step halves that uncertainty. If the residual at a returned state is 10^{-3} , the remaining potential error is at most $2 \cdot 10^{-3}$; the factor $1/(1-q)$ sums all the smaller corrections that can follow. When a sensor changes slightly, the preceding equilibrium is a useful starting point because Theorem A.4 bounds how far the required answer can move. This can save repeated work on a slowly changing stream. It gives little advantage after a large change, and convergence can be expensive when q is close to one. A unique equilibrium also erases differences in starting state: lasting history must enter through records, traces or a changed drive, rather than through the warm start alone.

Theorem A.5 (bounded damage). *Let $|\text{act}| \leq A$, let S be a set of neurons, let W^S be W with the columns in S set to zero, and let $\mu_S = \max_i \sum_{j \in S} |W_{ij}|$. If F_c satisfies Theorem A.1, then the lesioned map satisfies it with the same constant, and $\|v^* - v^{*,S}\|_\infty \leq dt \mu_S A / (1-q)$.*¹⁹

Theorem A.6 (the library's neuron). *For $\lambda \geq 0$ and $\ell \geq 0$, the activation (5) is Lipschitz with constant $\frac{\lambda}{4} \max(\frac{1}{1-\rho}, \frac{\ell}{\rho})$, satisfies $|\text{act}| \leq \max(1, \ell)$, and publishes exactly zero at rest. Under the learning model the constant is $1/2$ and the bound is 1, so every net whose row mass is below 2 settles by a certified contraction for every $0 < dt \leq 1$.*²⁰

¹⁵Lean: `settle_contracting`.

¹⁶Lean: `equilibrium_unique`, `silence`.

¹⁷Lean: `apriori_bound`, `aposteriori_bound`.

¹⁸Lean: `equilibrium_lipschitz_stimulus`, `warm_start_bound`.

¹⁹Lean: `lesion_bound`, `lesion_contracting`.

²⁰Lean: `activation_lipschitz`, `activation_abs_le`, `activation_zero`, `learningActivation_lipschitz`, `learning_model_certified`.

For the hypothesis, the certificate makes the numerical state of a changing brain checkable: when its conditions hold, a returned decision carries a bound on its distance from equilibrium. This bound certifies the solve. It does not certify the accuracy of the learned body model or the safety of an action.

A.4 Learning by two equilibria

The free phase settles under the stimulus alone to s^0 . The positive nudged phase settles from s^0 with an extra drive of strength β on the output neurons toward a target: for a quadratic nudge the drive is $\beta(\text{target}_i - s_i)$, for a cross-entropy nudge $\beta(\text{target}_i - p_i)$ with p the softmax of the outputs at temperature T , one softmax per declared slot. The centered estimator settles a negative nudged phase at $-\beta$ from the same free state. With s^+ and s^- the two nudged equilibria, the contrast at synapse (i, j) and at neuron i is

$$\Delta_{ij} = \frac{1}{B} \sum_b \frac{s_{b,i}^+ s_{b,j}^+ - s_{b,i}^- s_{b,j}^-}{2\beta}, \quad \Delta_i = \frac{1}{B} \sum_b \frac{s_{b,i}^+ - s_{b,i}^-}{2\beta}. \quad (6)$$

Reciprocal synapses share one parameter and receive the mean of their two contrasts; a cap bounds every efficacy.

Lemma A.7 (locality). *There is a fixed function g of four real numbers such that the contrast of synapse (i, j) equals $g(s_i^+, s_j^+, s_i^-, s_j^-)$; the contrast is symmetric in i and j ; and if the two compared equilibria agree, every contrast is zero.*²¹

The lemma describes the information used by the update. Reverse-mode differentiation retains or re-computes intermediate activations; the equilibrium contrast instead reads each synapse’s endpoints in the perturbed states. This can reduce storage when differentiating through many settling iterations. Both perturbed equilibria must be computed, and implicit differentiation offers another way to avoid retaining the iteration history [4].

Fact A.8 (the contrast is the gradient). *On a smooth stable equilibrium branch with symmetric effective recurrent weights, converged phases and no adaptation, the limit $\beta \rightarrow 0$ of the contrast is minus the gradient of the nudge’s loss with respect to the weight of a reciprocal pair [54]; the centered estimator cancels the first-order bias in β [31]. The library’s tests check the identity numerically: at $\beta = 10^{-4}$ and phase residuals below 10^{-11} the contrast equals the finite-difference gradient to a relative tolerance of 10^{-5} .*

Free and nudged phases are the deterministic form of the two phases of the Boltzmann machine [1, 36], equivalent to backpropagation in a layered net [65], and the local-activation-difference rule of generalized recirculation [45]. Reward uses the three-factor form: an eligibility trace $e \leftarrow \gamma \lambda e + \Delta$ of the same contrast, a linear critic with its own trace, and one broadcast error $\delta = r + \gamma V(s') - V(s)$, with the update $\eta \delta e$ [16, 57, 58].

Intuition: let a small correction travel through the circuit. The target gently pushes the outputs, and recurrence transmits that change to the rest of the net. A synapse compares how strongly its two endpoints coactivate under the positive and negative pushes. Their difference tells it which direction the connection should move under the assumptions of Fact A.8. Credit reaches the synapse through the circuit’s changed state, so the update itself needs only its endpoints. There is a cost: the learner must solve the perturbed equilibria, and a locally computed update can be slower than a conventional backward pass. An eligibility trace carries that local comparison forward when reward arrives later; it does not guarantee that one delayed scalar identifies the right cause. The advantage is the form and storage of the update, with learning quality established by the controlled comparisons.

B Optional record algebra and capacity

A record cortex keeps, for each of its cells, what followed the readings that cell took part in. A reading $x \in \mathbb{R}^d$ of declared source neurons is made mean-free by a running mean $\bar{x}$, projected by a fixed random matrix R onto $D \gg d$ cells, rectified, and reduced to the a strongest entries at unit length,

$$k(x) = \text{kwta}_a(\max(R(x - \bar{x}), 0)). \quad (7)$$

²¹Lean: `contrast_local`, `contrast_symm`, `contrast_eq_zero_of_agree`.

Each predicted field holds one table $M \in \mathbb{R}^{D \times m}$. The prediction is the read $k(x)M$, and the witnessed outcome y is written by the delta rule [27, 63] at rate η ,

$$M \leftarrow M + \eta k(x)^\top (y - k(x)M). \quad (8)$$

Consequence fields use a slow rate so that a regularity is averaged over many outcomes; valued fields such as reward use rate one, so that one exposure writes. The optional record store is distinct from the neural core and from the fixed-model Gaussian boundary.

Theorem B.1 (one-shot memory and its interference law). *For any record M , keys k , q , value y and rate η over a commutative ring: (i) the read at q after the write changes by $\eta(q \cdot k)$ times the correction $y - kM$; (ii) if $k \cdot k = 1$, the read at k after a write at rate one is exactly y ; (iii) if $q \cdot k = 0$, the read at q is unchanged; (iv) if $k \cdot k = 1$, writing the same correct observation twice at rate one changes nothing.*²²

Theorem B.2 (how far one write reaches). *Over the reals, the change a write at k makes to the read at q has sup norm exactly $|\eta| |q \cdot k|$ times the sup norm of the correction, and an overlap of at most ϵ bounds it by $|\eta| \epsilon$ times that norm.*²³

Theorem B.3 (exact storage and capacity). *After writing N pairwise orthogonal unit keys in sequence at rate one, every one of them reads its own value, whatever the initial record was; and over a field at most d keys can be pairwise orthogonal unit vectors in a d -dimensional key space.*²⁴

Theorem B.4 (pattern separation). *Two keys whose supports are disjoint are orthogonal, so a write at one leaves the read at the other exactly unchanged; a family of unit keys with pairwise disjoint supports, written in sequence at rate one, is stored exactly from any initial record; and a code of width D with a active cells admits at most $\lfloor D/a \rfloor$ such keys.*²⁵

Intuition: a correction follows the overlap of two addresses. Suppose a reading activates two cells equally, giving a unit key $k = (1, 1, 0, 0)/\sqrt{2}$. A write at rate one corrects that reading in a single step. A second reading with key $q = (0, 0, 1, 1)/\sqrt{2}$ shares no cells and keeps its prediction exactly. A third, $q = (0, 1, 1, 0)/\sqrt{2}$, shares one cell and receives half the first reading’s correction, since $q \cdot k = 1/2$. This is the interference law in a small example: the same overlap that permits generalization can overwrite another fact.

The four theorems explain why a useful code can learn quickly from a stream: each observation directly corrects the relevant prediction while disturbing other predictions in proportion to their key overlap. Sparsity alone guarantees neither distinct addresses nor retention; two sparse keys can be identical, and finite capacity limits exact separation. With an appropriate code, direct writes can need fewer observations than fitting the same associations in distributed weights. The record theorem does not by itself prove a replay requirement for another architecture. The published examples document measured uses of these records: a composer whose held-out accuracy is measured with the records writing online [R99], soft bodies that write records during a life and never inherit them [R100], and a Connect Four value patch whose school leaves the store empty: in three runs that wrote every schooled position into a store of 4,096 cells, the slow readout alone named the winner about three points more often on positions never seen and two points more often on positions witnessed [R101]. In the same example two positions that differ by one stone share 51 percent of their active cells, and a write to one moves the other 53 percent as far, in line with the interference law above; observing a position together with its neighbors, each anchored at its current value, moves it by 0.71 and then by 0.04 [R101].

Theorem B.5 (attention is one read of a record). *Let $M = \sum_t k_t^\top y_t$ be the additive record of N key/value pairs. The read at query q is $qM = \sum_t (q \cdot k_t) y_t$, the stored values weighted by the query/key dot products. If the keys are one-hot at distinct addresses, the read at a one-hot query is exactly the stored value, for every N up to the size of the address space.*²⁶

²²Lean: `read_write, read_write_self, read_write_orth, repeat_write_stable.`

²³Lean: `read_write_dist, read_write_dist_le.`

²⁴Lean: `read_writeUpTo, capacity_le.`

²⁵Lean: `dotProduct_eq_zero_of_disjoint_support, read_write_of_disjoint, read_writeUpTo_of_disjoint, disjoint_family_card_le.`

²⁶Lean: `read_hebb, one_hot_recall.`

The first sentence is unnormalized linear attention [3, 26, 55]; the softmax read of a transformer differs by an exponential kernel and a normalizer, and modern Hopfield networks show that it is one step of an energy network [28, 29, 47]. The residual write of equation (8) is the delta rule over a sequence [66]. The identity covers additive dot-product attention, not the general softmax operation. Unlike an additive write, the residual write corrects its current prediction.

C Wiring and a declared cost model

Theorem C.1 (the certificate under evolution). *Let W satisfy Theorem A.1 with row mass ρ . (i) Every W' with $|W'_{ij}| \leq |W_{ij}|$ for all i, j satisfies it with the same constant: pruning a synapse, masking a projection, lowering a gain and scaling by any $|g| \leq 1$ are instances. (ii) If V has row mass μ and $L(\rho + \mu) < 1$, then $W + V$ settles by a certified contraction with constant $1 - \text{dt}(1 - L(\rho + \mu))$: growth by new projections stays certified inside the row-mass budget. (iii) Every member of a population whose genomes share one row-mass bound settles by one certificate, whatever its individual wiring.*²⁷

The row-mass budget certifies convergence of admitted recurrent weights. It does not establish useful specialization or certify a changing nonlinear memory basin.

C.1 Cost

Fix the declared cost model. A patch net with E synapses that settles in k steps costs kE multiply-adds per decision, plus Dm for a record read and the same for a write, and holds n potentials and Dm numbers whatever has happened before. A transformer with ℓ layers of width d over a window of T tokens costs about $\ell(12d^2 + 2Td)$ multiply-adds per token and holds a key/value cache of $2\ell Td$ numbers.

Theorem C.2 (what grows and what does not). *Under that model, the arithmetic a transformer spends on a stream of T tokens is $\ell(12d^2T + dT(T - 1))$, which grows quadratically in T , and its cache grows without bound. The arithmetic a patch net spends is $T(kE + 2Dm)$ and its state is constant in T . For every choice of positive ℓ, d, k, E, D, m there is a horizon beyond which the patch net's total is the smaller of the two.*²⁸

The theorem is arithmetic about the declared model, not a claim that the two objects compute the same function. It states where the shapes differ: one of them re-reads its past, the other has written its past into a fixed store. Measured costs depend on the numerical solver, query accuracy and full controller.

Intuition: compression trades history for a bounded state. For a repeated body-control task, a learner may need the present position and a learned map from actions to consequences, rather than every observation that established the map. Records and traces carry that summary between decisions. Their fixed size explains the linear stream cost in the theorem, and also imposes a limit: they cannot preserve arbitrarily many independent facts exactly. The quadratic comparison assumes attention to an expanding history; a transformer with a fixed window also has linear total cost in stream length. Whether retaining the summary is better depends on what the task requires, the cost of finding the representation, the achieved error and the measured latency.

D Theorem map and the Lean library

Every Theorem, Lemma and Corollary in the paper names the Lean declaration that proves it. The public Cadence library includes the certified recurrence, activation, local contrast and wiring results, together with the representation, noise, basin, specialization and temporal-overlap boundaries [37]. The optional linear-record algebra has a separate import; it is not a component of **PatchNet**. These sources pin Lean 4.29.1 and Mathlib at the matching release. Their axiom audit allows only propositional extensionality, choice and quotient soundness, with no admissions or project-specific axioms. The historical cost results are retained in the

²⁷Lean: `rowMassLE_of_abs_le`, `rowMassLE_smul`, `rowMassLE_add`, `mutation_contracting`, `growth_contracting`, `population_contracting`.

²⁸Lean: `attention_total`, `attention_total_ge_quadratic`, `record_total`, `record_state_const`, `exists_crossover`.

full manuscript’s proof collection; exploratory structural-memory, readable-trace, causal-path and quadratic-boundary results have separate sources and verification in the accompanying evidence [38]. The causal-path, quadratic-boundary and protected-path lemmas are included in the public package’s 169 audited statements; other exploratory collections remain separate. The full manuscript build checks these additional sources as well. None of these collections is a formal refinement proof of the floating-point experiments or the complete library interface.

Table 9: The theorem map. Settling theorems hold for Lipschitz activations; the Activation module instantiates them for this neuron. Memory identities hold over commutative rings, the capacity bound over fields, the cost statements over the naturals, the depth identities over real matrices.

paper	module and declaration	content
Thm A.1	Settling, <code>settle_contracting</code>	contraction with constant $1 - dt(1 - L\rho_W)$
Cor A.2	Settling, <code>equilibrium_unique</code> , <code>silence</code>	any fixed point is the equilibrium; rest is the equilibrium of no input
Thm A.3	Settling, <code>apriori_bound</code> , <code>aposteriori_bound</code>	the two error certificates
Thm A.4	Settling, <code>equilibrium_lipschitz_stimulus</code> , <code>warm_start_bound</code>	the equilibrium moves by at most $dt\delta/(1 - q)$; k warm steps leave $dt\delta q^k/(1 - q)$
Thm A.5	Settling, <code>lesion_bound</code> , <code>lesion_contracting</code>	a lesion moves the equilibrium by at most $dt\mu_S A/(1 - q)$
Thm A.6	Activation, <code>activation_lipschitz</code> , <code>activation_abs_le</code> , <code>activation_zero</code> , <code>learningActivation_lipschitz</code> , <code>learning_model_certified</code>	the neuron’s Lipschitz constant, range and rest; row mass below two certifies the learning model
Lem A.7	Credit, <code>contrast_local</code> , <code>contrast_symm</code> , <code>contrast_eq_zero_of_agree</code>	locality, symmetry, zero contrast for identical states
Thm B.1	Memory, <code>read_write</code> , <code>read_write_self</code> , <code>read_write_orth</code> , <code>repeat_write_stable</code>	interference law, one-shot write, orthogonal invariance, idempotence
Thm B.2	Sparse, <code>read_write_dist</code> , <code>read_write_dist_le</code>	the exact size of one write’s effect on another read
Thm B.3	Memory, <code>read_writeUpTo</code> , <code>capacity_le</code>	exact storage of an orthonormal family; at most d such keys
Thm B.4	Sparse, <code>dotProduct_eq_zero_of_disjoint_support</code> , <code>read_write_of_disjoint</code> , <code>read_writeUpTo_of_disjoint</code> , <code>disjoint_family_card_le</code>	disjoint supports never interfere; a disjoint family is stored exactly; a code of width D with k winners holds at most $\lfloor D/k \rfloor$ of them
Thm B.5	Memory, <code>read_hebb</code> , <code>one_hot_recall</code>	the additive read is linear attention; one-hot recall is exact
Thm C.1	Wiring, <code>rowMassLE_of_abs_le</code> , <code>rowMassLE_smul</code> , <code>rowMassLE_add</code> , <code>mutation_contracting</code> , <code>growth_contracting</code> , <code>population_contracting</code>	pruning, masking and rescaling keep the certificate with the same constant; growth keeps it inside the row-mass budget; one bound certifies a whole population
Thm C.2	Cost, <code>attention_total</code> , <code>attention_total_ge_quadratic</code> , <code>record_total</code> , <code>record_state_const</code> , <code>exists_crossover</code>	the transformer’s arithmetic over a stream is quadratic in its length and its cache unbounded; the patch net’s is linear with constant state; a crossover exists
Thm 2.2	Depth, <code>readback_collapse</code> , <code>tower_eq_mulVec</code> , <code>drive_eq_collapsed</code> , <code>equilibrium_tower_iff</code> , <code>settle_tower_eq</code>	a tower of linear readbacks is one matrix; its equilibria and settling trajectories are those of a depth-zero patch with the collapsed weights

Three further results are used as facts and are Lean-checked in the observer-patch corpus rather than in this package [40, 42]: local repair to consensus on a federation of patches computes every finite Boolean function (`echosahedral_universality`); a work-conserving scheduler on a ranked repair system terminates

without a fairness premise; and the accepted repair steps of a fixed input-independent federation number at most $n(n+1)/2$, order-sharp. No embedding of those discrete programs into this neuron model is supplied.

E Evidence map

Empirical numbers are read from the named receipts copied into the evidence package, with each receipt’s own digest where it has one, its implementation version or source hashes, and the sha256 of the copied bytes recorded in the package manifest. Quoted external results have bibliographic citations. Exact mathematical constants follow their stated definitions and hypotheses. The memory-proof sources are copied under the manifest label `mission-memory-proofs`.

The causal-admission and partitioned-phrase packets include numerical observations, checkpoints, archived implementations and verification receipts. Source hashes identify the recordings used for their audio checks; raw recordings and rendered audio are excluded from these packets. Numerical replay and reproduction of the audio from its identified source therefore have different input requirements. Deterministically compressed copies retain the original source-byte hashes in the manifest.

Five examples are published with their sources, pages and checks at <https://github.com/muellerberndt/cadence-examples>; four are pinned here at one commit and the dozing cat at a later one. A worm whose brain is the measured *C. elegans* connectome takes a lesson whenever food or pain arrives and drives its body from its command interneurons; its receipt holds the browser brain and body to the library and to the reference body, and a receipt of conditioning with controls is work in progress [R98]. A composer starts from silence and computes a track, hearing each half-beat it plays [R99]. Soft bodies inherit their shape, muscles and cortices, learn during a life and plan their motor through a learned model while every unit of computation costs mass [R100]. A Connect Four brain plans by a supplied search over imagined boards that reads a value learned from 4.4 million positions of a perfect player’s games; in 200 games graded move by move it wins or draws nineteen of twenty going first against that player and loses none to an AlphaZero baseline [R101]. Each states what is supplied, what is learned, what was measured and what it does not show. The dozing cat runs the whole brain of Section 5.1 in the browser with the library’s arithmetic at <https://floatingpragma.io/cadence-examples/dozing-cat/>, held to the library by a parity test on a committed recording [R117]. The other acceptance demos, the room, the arm and the habit experiments are held with their receipts and reports in the evidence package under the demo and habit labels of Table 10.

Table 10: Receipt identifiers and manifest labels. The library column names the release whose sources the receipt binds; OPN is the earlier observer-patch-nets library. The manifest preserves source references and copied-file hashes.

id	manifest label	receipt	library
R1	recall	associative recall over 128 pairs, two transformer arms	0.5.0
R3	pattern-separation	correlated keys under revision streams, six arms, five seeds	0.9.0
R5	digit-span	digit span, six arms, fade sweep	0.7.0
R6	n-back	<i>n</i> -back, four arms, two baselines	0.7.0
R8	adaptive-depth	compute follows change: warm steps, budget, bounds, three row masses, five seeds	0.9.0
R10	event-repair	event-driven repair on the fly graph, 135 conditions	0.8.1
R12	surprise	settling steps against entropy	0.5.0
R13	matched-learning	three learning rules on one 48-neuron mechanism	0.8.1
R16	lesion	damage and recovery, three seeds	0.5.0
R17	permuted-mnist	thirty permuted tasks, six arms	0.5.0
R22	fly-counterfactuals	whole-fly counterfactual screening	0.8.1
R23	opn	fly wiring against shuffled, conformance	OPN
R24	mouse	mouse labyrinth, connectome against shuffled and a network	0.7
R26	player	self-play of the many-games brain	0.7
R27	player	sixteen games at once with a dopamine floor	0.8
R31	sizing	text width sweep, per-run files	0.4 to 0.7
R32	author	language pilots at level one	0.7
R33	hopper, pong	control gates, logs	0.5.0
R34	repeat-previous, slippery-ant	partially observed cards; slippery ant	0.7 to 0.8.1
R36	split-mnist	split MNIST, one head	0.5.0
R37	worm-probe	supplied worm mechanism under lesions	0.8.1

id	manifest label	receipt	library
R41	ep-inputs	fixed-input derivatives and free/free asymmetry control	0.9.0
R42	rehearsal	split and permuted MNIST at disclosed memory and work budgets	0.9.0
R43	content-prototypes	noisy cues, fixed prototypes, aliases and capacity	0.9.0
R44	policy-credit	delayed choices, loss-policy correction and score control	0.9.0
R45	sequence-readback	causal held-out text and validation grids	0.9.0
R46	drift-diagnostic	thirty tasks and equal-budget reset forks	0.9.0
R57	instrument	the instrument: five acceptance seeds, ten gates	0.9.0
R58	composer-stage	the chip composer: five acceptance seeds, thirteen gated and seven open measures	0.9.0
R61	wiring-search	wiring search on four families and on the composer’s streams: evolution, random search, hand-found wirings, intervention states	0.9.0
R62	scaling-curves	fourteen arms on four task families at four budgets, preregistered floors, retention read-back	0.9.0
R63	combined-learner	banks beside a gradient patch with a guard, ten arms, four budgets	0.9.0
R64	evolved-development	the evolution loop over wiring and development policy: contract and calibration	0.9.0
R65	composer-curves	the chip composer as curves over a life: three wirings, two sizes, two life lengths, four seeds	0.9.0
R66	player-curves	the 1942 records controller and a PPO convolutional net at matched frames, three seeds	0.9.0
R69	player-1943	the 1943 imitation player and its practice on 1942, frozen tests on eight starts	0.9.0
R71	mission-basin	nonlinear cue persistence, opposing evidence and continued-learning stress; source, full receipt and independent verification	OPN
R72	mission-specialization	nine-patch completion and warm revision; source, full receipt and independent verification	diagnostic
R73	mission-structural	normalized-contact repair, continuous controls and shared-connection failure; full receipt and verification	diagnostic
R74	mission-active-evidence	exact finite selective-inspection reference; frozen model, exhaustive enumerator and receipt	reference
R75	library-memory	reciprocal library core without temporal boundary; three-seed acquisition, retention, cue and continuation checks with archived sources	core
R76	library-temporal-boundary	same declared library tasks with held temporal overlap; finite context, cold-acquisition failures and archived sources	core
R77	amen-patch-features	fixed-seed recorded-audio feature prediction with and without temporal overlap; matched autoregressive controls, arrays and source snapshots	core
R78	library-selective-retention	absent-boundary, fixed-input and calibrated selective-plasticity tests; complete outcomes, checkpoints, sources and independent endpoint verification	core
R79	scalar-continuation	periodic continuation with current, temporal, explicit-history and supplied-clock inputs; frozen protocol, full receipt and arithmetic verification	core
R80	readable-trace	six scalar trace-repair lemmas; pinned sources and compiled axiom audit	formal
R81	intent-control	brief-cue motor/consequence test, matched trace ablation, private branches, continued-state acquisition and linear-decoder diagnostics	core
R82	intent-episodic	independent episodic intentions with fresh-state training; trace ablation and supplied intention/position acquisition control	core
R83	temporal-repair	causal-path uniqueness proof; linear and nonlinear temporal credit, initialization comparison and fresh motif validation	diagnostic
R84	temporal-actor	linear body acquisition by centered detuning; joint action planning, actual feedback, cold restoration and independent execution audits	diagnostic
R85	temporal-observer	noisy past-state repair and one-record Gaussian-prefix equivalence; actual records, model binding, controls and independent all-history audits	diagnostic
R86	quadratic-boundary	six scalar fixed-model quadratic elimination and observation-admission theorems; compiled axiom audit	formal
R87	revision-exposure	fresh crossed static/warm exposure and gradient/Newton repair; complete memory outcomes and independent endpoint arithmetic	diagnostic
R88	amen-temporal-capacity	real-audio intention-conditioned finite paths; all-seed acquisition, retention controls and independent endpoint/recurrence checks	capacity
R89	amen-temporal-retention	local response constraints, new-phrase learning, protected-path Lean induction, sample listening and independent checks	conditional
R90	amen-sample-events	two recorded-bar recalls, supplied sample actions, known-action controls and independent neural/audio replay	capacity

id	manifest label	receipt	library
R91	amen-sample-retention	sampled-action retention, protected-path conditioning, matched control and independent neural/PCM checks	bound
R92	nonlinear-input-planning	continuous input detuning, acquired nonlinear-body control, full training/plan replay and actual readback controls	conditional
R93	amen-sample-repertoire	joint four-bar recorded repertoire, all taught sample identities, untaught cue interpolation and exact audio replay	capacity
R94	amen-sample-metric	matched constrained-readout metric, complete update and PCM replay, conditioning and finite-detuning counterexample	conditional
R95	causal-admission	observed-loss parameter admission, acquired acoustic control and fixed-reference restoration; bounded outcomes and independent replay	conditional
R96	partitioned-cue	masked joint recurrence, capped dense comparison, complete update replay and untrained radius-normalization controls	diagnostic
R97	self-readback-control	exact matrix readback equivalence and finite temporal sensitivity controls	diagnostic
R98	example-worm	published example: the <i>C. elegans</i> connectome as one masked temporal patch net; browser brain and body held to the library and to the reference body	main
R99	example-amen	published example: a composer in one record patch; next-event accuracy on held-out tracks against baselines, browser engine held to the library	main
R100	example-patchworld	published example: soft bodies with two record patches under selection; conserved mass, the founders' model on a held-out stream, two base worlds and six worlds with the night and the sleep gene, 20,000 ticks each	main
R101	example-connect4	published example: a value patch schooled by a perfect player and a supplied search; 200 games with every move graded by the solver, the search alone as control, the record store under bulk writing, and what neighboring positions share in it	main
R102	language-probes	order, agreement across attractor nouns and recitation probes for symbol streams; one patch, two stacked patches, a GRU and a transformer	twin
R103	language-records	a grammar from records alone: one and two passes of delta-rule writes, held-out combinations, minimal pairs	twin
R104	language-scaling	held-out loss and validity against width and the share of taught pairs at equal updates, three learners	twin
R105	language-speaker	the two-layer patch, a GRU and a transformer trained by cross-entropy on the same split; recitation of the read sentences	twin
R106	language-one-shot	sentences never seen by the slow weights written once into the store; retention and the disturbance of the library	twin
R107	language-settling	a settled-sentence speaker on the equilibrium net: calibration, slot accuracy after one epoch, unconverged settles	main
R108	language-sleep	sleep in Wagner's design: the day's store, the awake control with the teacher, dreams from the day's cues and from recombined cues, the teacher for the same updates	twin
R109	example-connect4-sleep	the value patch under the day and night regime: one night from chance at three store sizes, matched update counts against the school, an averaging store, one day and three nights, twenty games an opponent graded by the solver	main
R110	example-amen-sleep	the composer under the day and night regime: held-out next-event accuracy at bedtime and after each night under every record variant, against three epochs with the corpus	main
R111	example-patchworld-sleep	a sleep gene in the soft-body world: three worlds of 20,000 ticks with two seeds, the share of sleepers over time, lifetimes and late-life model error by sleep	main
R112	language-config	the record store's configuration one knob at a time on the grammar: cells, active cells, output code width, write rate, passes, context width, timescales and address weights	twin
R113	language-correction	sleep that keeps only sure dreams, and dawn re-reading of the unsure cues, against the sleep and teacher arms	twin
R114	demo-shell-game	the shell game: the belief patch, the record patch and the tracker on five worlds, by stretch length; the genome against the hand-set genome and a random search; timing; the page driven	0.13.0
R115	demo-eye-and-ear	record patches joined by ports: two coupled streams, six arms, five seeds; the ball with the lights out and the room muted, the port cut; the topology as a genome on both; single-process timing	0.13.0

id	manifest label	receipt	library
R116	demo-doing-cat	the doing cat: four arms on five lives, the governor patch's and the threshold rule's genomes under selection against random search, held-out lives, timing, the page driven	0.13.0
R117	example-doing-cat	the published cat example: its receipts, the parity recording of the library's decisions that the browser brain is held to, and its card	0.13.0
R118	demo-ventriloquist	the ventriloquist: four arms and two ablations on five worlds with the readback per phase, the steering patch's genome under selection against random search with cut and shuffled finals, the page driven	0.13.0
R119	demo-room	the room: governed against never learning, always imagining and always learning on the window and the heater, the four-way decomposition, the genome under a compute price and under a surprise price, quiet-machine timing; Pong with the rules swapped mid-life	0.13.0
R120	habit-cartpole	habits as patches on the cart-pole: distillation from the gains, the matched categorical coding and reward in imagination on ten seeds, the tilt and its fix loop, the genome under selection on three seeds	0.13.0
R121	demo-arm	the reaching arm: the body's diagnostics, the driven run with the weight for the hand-designed and the evolved genome, the habit's genome under selection	0.13.0
R122	habit-pong	the Pong habit: the planner's corpus, the habit and the linear readout trained and played on ten episodes, the second round on the habit's own situations	0.13.0

F Notation

n neuron count; v_i potential; $s_i = \text{act}(v_i)$ activation; W_{ij} effective weight from j to i ; c_i stimulus including bias; dt step; ρ_W maximum absolute incoming weight sum; L activation Lipschitz constant; q contraction constant; v^* equilibrium; β nudge strength; $s^\pm$ nudged equilibria; Δ contrast; η learning rate; M record table; k, q keys; y value; D code width; a active code cells; $k(x)$ reading code; R code expansion; e eligibility trace; δ reward error or stimulus change as stated; γ, λ discount and trace decay; μ added-projection row mass; ℓ, d, T transformer layers, width and window in the cost model.